

]

TrackVision

Real-time Client-Side Multiple Object Tracking in Video

Unit 2 Computer Vision Mini Project

Submitted in partial fulfillment of the requirements for the degree of

Bachelor of Technology in Computer Science and Engineering

Student:	[Student Name]
Roll Number:	[Roll Number]
Branch:	Computer Science and Engineering
Semester:	[Semester]
Academic Year:	2025–2026
Submission Date:	12 September 2026
Guide:	[Guide Name]
Department:	Department of Computer Science and Engineering
Institute:	[Institute Name]

Certificate

This is to certify that the project report entitled “**TrackVision: Real-time Client-Side Multiple Object Tracking in Video**” submitted by **[Student Name]** (Roll No: [Roll Number]) in partial fulfillment of the requirements for the award of the degree of **Bachelor of Technology in Computer Science and Engineering** is a bonafide record of work carried out by him/her under my supervision and guidance.

The project represents the student’s own work and has not been submitted elsewhere for any other degree or diploma.

[Guide Name]

Project Guide

Department of CSE

Date: August 24, 2026

[Head of Department]

Department of CSE

[Institute Name] stitute Name>

stitute Name

Date: August 24, 2026

Declaration

I, **[Student Name]**, Roll Number **[Roll Number]**, student of **Bachelor of Technology in Computer Science and Engineering**, hereby declare that the project report entitled “**TrackVision: Real-time Client-Side Multiple Object Tracking in Video**” is an original work carried out by me under the guidance of **[Guide Name]**, Department of Computer Science and Engineering, **[Institute Name]**.

This work has not been submitted to any other university or institution for the award of any degree or diploma. All sources of information used in this report have been duly acknowledged.

[Student Name]

Roll No: **[Roll Number]**

Date: August 24, 2026

Place: **[City]**

Abstract

Multiple Object Tracking (MOT) is a fundamental computer vision task that extends object detection by maintaining consistent object identities across consecutive video frames. Traditional MOT systems typically operate on powerful servers with GPU acceleration, requiring video uploads and introducing latency, privacy concerns, and bandwidth costs.

This project presents **TRACKVISION**, a fully client-side MOT system that runs entirely within a web browser using **ONNX Runtime Web** for hardware-accelerated inference. The system implements a custom **BYTETRACK++** algorithm with a full 6x6 covariance Kalman filter, Hungarian data association, and optional OSNet x1.0 appearance-based re-identification. Two detection modes are supported: a fast path using YOLOv8n (COCO-80 classes, 12.8 MB) and an open-vocabulary path using YOLO-World with CLIP ViT-B/32 text embeddings computed in-browser.

The architecture employs four Web Workers (YOLO detection, YOLO-World, ReID, Tracker) communicating via a typed message protocol, with the main thread handling class-aware NMS, visualization, and state management via Zustand. Video input is captured via **OffscreenCanvas** at 640x640 letterboxed resolution. The system achieves sub-millisecond JavaScript-side tracking overhead (0.31–1.04 ms/frame at 5–20 detections) and provides real-time telemetry (FPS, latency breakdown, execution provider).

Key contributions: (1) A production-grade client-side MOT pipeline with zero server dependencies, (2) Dual detection modes with in-browser CLIP text encoding, (3) Custom TypeScript ByteTrack++ with full covariance Kalman filtering and appearance fusion, (4) Comprehensive UI with Ghost/Follow/Replay modes, 5-minute history, and analytics, (5) 49 unit tests validating core algorithms, (6) Static deployment with COOP/COEP headers for WebGPU/SharedArrayBuffer support.

Limitations: No formal benchmark dataset is used; standardized MOT metrics (MOTA, IDF1, HOTA) are not computed. Evaluation consists of implementation-level unit tests, microbenchmarks, and qualitative observation. Future work includes MOTChallenge evaluation integration and video file upload support.

Keywords: Multiple Object Tracking, ByteTrack, YOLOv8, YOLO-World, Kalman Filter, Hungarian Algorithm, OSNet, CLIP, ONNX Runtime Web, WebGPU, Client-Side Inference

Acknowledgement

I would like to express my sincere gratitude to all those who have contributed to the successful completion of this project.

First and foremost, I am deeply grateful to my project guide **[Guide Name]** for their invaluable guidance, constant encouragement, and constructive feedback throughout the project. Their expertise and insights were instrumental in shaping the technical direction of this work.

I extend my thanks to **[Head of Department]**, Head of the Department of Computer Science and Engineering, for providing the necessary infrastructure and support.

I am thankful to the faculty members of the Department of Computer Science and Engineering for their academic guidance and for creating an environment conducive to learning and innovation.

This project stands on the shoulders of giants in the computer vision and open-source communities. I acknowledge the authors and maintainers of:

- **ByteTrack** – YOLOX/ByteTrack [1]
- **YOLOv8** – Ultralytics [2]
- **YOLO-World** – AILab-CVC/YOLO-World [3]
- **OSNet** – Kaiyang Zhou [4]
- **CLIP** – OpenAI [5]
- **ONNX Runtime Web** – Microsoft [6]
- **React, Vite, Tailwind CSS, Zustand, Recharts** – Respective open-source communities

Special thanks to the developers of `onnxruntime-web`, `wkentaro/yolo-world-onnx`, and `Xenova/transformers.js` for providing the ONNX models and tokenizer assets that make in-browser open-vocabulary detection possible.

Finally, I thank my peers and family for their unwavering support and patience during the long hours of development and documentation.

[Student Name]

Roll No: [Roll Number]

Contents

1	Aim	12
2	Introduction	14
3	Literature Review	18
4	Dataset Description / Details	24
5	Methodology	27
6	Algorithm / Flowchart	34
6.1	Computational Complexity	39
7	Source Code	40
8	Output Screenshots	46
9	Result & Accuracy	53
10	Applications	60
11	Conclusion	63

chapterList of Figures

8.1	TrackVision Landing Page – Hero section, Feature Grid, Open Vision Playground, Performance Comparison, Use Cases, Code Quickstart, FAQ, and CTA Banner. Full-page scroll capture.	47
8.2	Command Center – Initial state after clicking “Launch Command Center”. Shows top navigation (Live, Tracks, Vision, History, Scene, Analytics), center viewport with “Start Tracking” button, and Engine Thinking Stream pill. Models not yet initialized. . .	48
8.3	Command Center – Ready state after model initialization (YOLOv8n loaded). “Start Tracking” button enabled, telemetry stream active, Live badge shows standby FPS.	48
8.4	Vision Panel (Detection Engine) – Fast Mode (YOLOv8n, COCO-80) vs Open Mode (YOLO-World + CLIP) selector. Open mode shows concepts input textarea with comma-separated prompts. Runtime fallback status shows execution provider.	48
8.5	Analytics Panel (Settings tab) – Real-time telemetry cards (FPS, frame time, inference latency, tracking latency), object volume AreaChart, class distribution BarChart. Updates live during tracking. . . .	48
8.6	Track Inspector (Tracks tab) – Active tracks list with track ID, class label, confidence score, status (NEW/TRACKED/LOST), and selection controls. Click to enable Follow mode.	49
8.7	Scene Map (Scene tab) – 640x480 top-down spatial density map with live track position dots, color-coded by track ID. Updates in real-time during tracking.	49

8.8	Timeline (History tab) – Gantt-chart style horizontal bars per track ID, 5-minute sliding window, variable playback speeds (0.5x, 1x, 2x, 4x), scrubber for frame-perfect seek.	49
8.9	Command Palette – Activated via Ctrl+K. Quick actions: Start/Stop Tracking, Toggle Ghost Mode, Toggle Follow Mode, Clear History, Switch Vision Mode, Open Settings.	49
8.10	Mobile Responsive View (375x667) – Bottom navigation bar (Live, Tracks, History, Scene, Vision, Analytics), stacked panels, touch-optimized controls.	50
8.11	Full Desktop Command Center – Live view with camera placeholder, left sidebar (Tracks panel shown), top navigation with telemetry (FPS, latency), Engine Thinking Stream, corner controls (Ghost Mode, Follow Mode), bottom live HUD.	51
8.12	Early detection frame – YOLOv8n output showing bounding boxes, COCO class labels, and confidence scores. No tracking IDs assigned yet (first frame). .	51
8.13	Three or more objects tracked simultaneously with persistent integer IDs, class labels, and confidence scores. IDs remain stable across frames.	51
8.14	Identity continuity demonstration – same physical object maintains the same track ID across multiple frames despite motion and minor appearance changes.	51
8.15	Ghost mode enabled – 15-frame fading trajectory trails with opacity decay, showing historical path of each tracked object.	51
8.16	Follow mode enabled – camera viewport locked to selected track, background dimmed, selected track highlighted with persistent ID and telemetry.	51

8.17	Object temporarily occluded – track maintains ID via Kalman prediction (interpolation frames), reappears with same ID after occlusion.	51
8.18	Multiple overlapping objects in close proximity – distinct IDs maintained, track merging activates only when $\text{IoU} > 0.5$, same class, and appearance similarity > 0.7	51
8.19	Track re-acquired after prolonged occlusion – OSNet appearance embedding fusion enables correct ID assignment despite spatial gap. Requires OSNet model loaded.	51
8.20	Real-time Telemetry – FPS history (line chart), latency breakdown (stacked area: preprocess, inference, postprocess, tracking, render), object count per frame, class distribution bar chart. Updates every frame during tracking.	52
9.1	Tracking pipeline latency scales sub-linearly with detection count. Model inference (not shown) dominates end-to-end frame time.	55

chapterList of Tables

3.1	Comparison of MOT methods relevant to this work. Bold indicates methods implemented in TrackVision.	22
4.1	Input video parameter handling	25
4.2	Dataset description summary	26
6.1	Per-frame computational complexity (asymptotic) .	39
8.1	Screenshot capture summary	46
9.1	Unit test coverage by module	54
9.2	Tracking pipeline latency (JS only, no model inference)	55
9.3	ONNX session initialization time (model load + compile)	55
9.4	Production bundle sizes	56
9.5	Recommended MOT metrics for future benchmark evaluation	57
9.6	Summary of measured results (achieved)	58
10.1	TrackVision deployment advantages for real-world applications	61

chapterAim

The aim of this project is to design, implement, and evaluate a real-time, fully client-side multiple object tracking system capable of detecting, associating, and maintaining persistent object identities across consecutive video frames directly within a web browser, without requiring any server-side computation or video upload.

The specific objectives are:

1. Client-side inference pipeline: Implement a complete MOT pipeline (detection $\rightarrow$ NMS $\rightarrow$ ReID $\rightarrow$ tracking $\rightarrow$ visualization) using ONNX Runtime Web with hardware acceleration (WebNN $\rightarrow$ WebGPU $\rightarrow$ WebGL $\rightarrow$ WASM cascade), ensuring zero data leaves the user's device.
2. Dual detection modes: Support both a fast, fixed-vocabulary mode (YOLOv8n, 80 COCO classes) and an open-vocabulary mode (YOLO-World + CLIP ViT-B/32 text embeddings) with in-browser tokenization and text encoding.
3. Custom ByteTrack++ tracker: Implement a tracking-by-detection algorithm with:
 - Full 6x6 covariance Kalman filter (constant acceleration model)
 - Hungarian algorithm for optimal data association
 - Track lifecycle management (TENTATIVE $\rightarrow$ CONFIRMED $\rightarrow$ LOST $\rightarrow$ REMOVED)
 - Appearance-based re-identification via OSNet x1.0 (512-dim embeddings)
 - EMA appearance fusion ($\alpha = 0.3$) and track merging for split-identity recovery
 - EMA bounding-box smoothing ($\alpha = 0.7$) and occlusion interpolation (up to 5 frames)

4. **Browser-native video acquisition:** Capture live camera feed via `getUserMedia()` and process frames using `OffscreenCanvas` at 640x640 letterboxed resolution with zero-copy buffer transfer to workers.
5. **Real-time telemetry and history:** Provide per-frame latency breakdown (preprocess, inference, postprocess, tracking, render), FPS monitoring, 5-minute rolling history with frame-perfect replay, and persistent track metadata (trajectories, distance, speed, bearing).
6. **Comprehensive UI/UX:** Implement a responsive Command Center with Ghost mode (trajectory trails), Follow mode (target lock), Time Machine (history scrubber), Scene Map (2D density), Analytics (class distribution, volume charts), and Command Palette.
7. **Validation and deployment:** Achieve 49 passing unit tests covering Kalman filtering, Hungarian matching, ByteTrack lifecycle, appearance fusion, NMS, and store history; produce a production build deployable to static hosts (Netlify, Vercel, Cloudflare Pages) with required COOP/COEP headers.

The system is evaluated through implementation-level tests, runtime microbenchmarks, telemetry observation, and qualitative tracking behavior assessment. Standardized MOT benchmark metrics (MOTA, IDF1, HOTA) are not computed in the current implementation due to the absence of a formal annotated dataset; this is documented as a limitation and future work direction.

Scope: The project demonstrates that a complete, production-quality MOT system can run entirely in-browser using modern web APIs and ONNX Runtime Web, achieving real-time performance on consumer hardware while preserving user privacy and eliminating server costs.

chapterIntroduction

section[Computer Vision]Computer Vision

Computer Vision (CV) is an interdisciplinary field that enables computers to derive meaningful understanding from digital images or videos. It seeks to automate tasks that the human visual system performs effortlessly: identifying objects, recognizing scenes, detecting motion, and understanding spatial relationships. Over the past decade, deep learning – particularly Convolutional Neural Networks (CNNs) – has revolutionized CV, achieving superhuman performance on benchmarks such as ImageNet classification [7], COCO object detection [8], and KITTI autonomous driving [9].

section[Video Analysis]Video Analysis

Video analysis extends CV from static images to temporal sequences. A video is a sequence of frames, typically captured at 24–60 FPS, where each frame is correlated with its neighbors. The temporal dimension introduces both challenges (motion blur, occlusion, appearance change) and opportunities (motion cues, track continuity, action recognition). Core video analysis tasks include action recognition, video object segmentation, and multiple object tracking.

section[Object Detection]Object Detection

Object detection answers the question: “*What objects are present in this frame, and where?*” Given an image, a detector outputs a set of bounding boxes with class labels and confidence scores.

Modern detectors fall into two families:

- Two-stage detectors (R-CNN [10], Faster R-CNN [11]): Region proposal → classification/regression. High accuracy, slower.
- Single-stage detectors (YOLO [12], SSD [13], RetinaNet [14]): Direct prediction from dense grid. Faster, suitable for real-time.

YOLO (You Only Look Once) pioneered real-time detection by framing it as a single regression problem. YOLOv8 [2] represents the current state-of-the-art in the YOLO lineage, offering improved accuracy-speed tradeoffs and a unified architecture for detection, segmentation, and pose estimation.

section[Object Tracking]Object Tracking

Object tracking answers: “*Where did the object detected in the previous frame go?*” Unlike detection (per-frame, independent), tracking maintains *identity* across time.

subsection[Single Object Tracking (SOT)]Single Object Tracking (SOT) SOT initializes with a ground-truth bounding box in frame 1 and predicts the target’s location in subsequent frames. Approaches include correlation filters (MOSSE [15], CSRT [16]) and Siamese networks (SiamFC [17], SiamRPN [18]). SOT assumes the target is always present and does not handle multiple objects or identity assignment.

subsection[Multiple Object Tracking (MOT)]Multiple Object Tracking (MOT) MOT extends SOT to *multiple simultaneous objects* with *unknown and variable count*. The core challenge is data association: given detections in frame t and tracks from frame $t - 1$, which detection belongs to which track? MOT systems typically follow the tracking-by-detection paradigm:

1. Detect objects in each frame independently.
2. Associate detections across frames to form trajectories.
3. Manage track lifecycle (initialization, confirmation, loss, deletion).

The distinction is fundamental:

Object detection answers: “What objects are present in *this* frame?”

Multiple object tracking answers: “Which detected object corresponds to *which previously observed object?*”

section[Problem Statement]Problem Statement

Despite advances in MOT algorithms, deployment remains a barrier. Most MOT systems:

- Require server-grade GPUs (CUDA, cuDNN, TensorRT).
- Necessitate video upload (bandwidth, latency, privacy).
- Depend on complex Python/C++ stacks (PyTorch, TensorFlow, OpenCV).
- Lack cross-platform accessibility (Windows/Linux/macOS, mobile).

There is a need for a zero-installation, privacy-preserving, cross-platform MOT system that runs entirely on the client device using only standard web APIs.

section[Motivation]Motivation

1. Privacy: Video data never leaves the device. Critical for surveillance, healthcare, retail analytics.
2. Latency: No network round-trip. Enables real-time interactive applications (AR, robotics, gaming).
3. Accessibility: Runs in any modern browser (Chrome, Edge, Firefox, Safari) on desktop or mobile.
4. Cost: Zero server infrastructure. Scales to millions of users with static hosting.
5. WebGPU: Emerging standard exposes GPU compute to the web, enabling near-native inference speed.

section[Objectives]Objectives

1. Implement a client-side MOT pipeline with YOLOv8n (fast) and YOLO-World + CLIP (open-vocabulary) detection.
2. Develop a custom ByteTrack++ tracker in TypeScript with full covariance Kalman filtering and Hungarian association.

3. Integrate OSNet x1.0 ReID for appearance-based occlusion recovery.
4. Build a responsive web UI with real-time telemetry, history replay, and multiple visualization modes.
5. Validate core algorithms with 49 unit tests; measure pipeline performance.
6. Deploy as a static PWA with WebGPU acceleration and HTTPS support for mobile camera access.

section[Scope]Scope

- In scope: Real-time browser MOT, dual detection modes, custom tracker, ReID, UI/UX, deployment.
- Out of scope: Standardized MOT benchmark evaluation (MOTA/IDF1/H video file upload (camera only), server-side components, 3D tracking, pose estimation, action recognition.

section[Project Contribution]Project Contribution

This project contributes a complete, production-grade, client-side MOT system implemented in TypeScript/React that demonstrates:

- The feasibility of running YOLO + ByteTrack + Kalman + ReID entirely in-browser via ONNX Runtime Web.
- A novel open-vocabulary detection path using in-browser CLIP text encoding (no external embedding service).
- A clean, typed Web Worker architecture with typed message protocols.
- Comprehensive UI/UX for MOT visualization and analysis.
- Open-source deployment model (static hosting, Docker, MIT license).

chapterLiterature Review

section[Traditional Object Tracking]Traditional Object Tracking

Before deep learning, tracking relied on hand-crafted features and probabilistic filters. The Kalman Filter [19] provides optimal state estimation for linear Gaussian systems. For visual tracking, the Mean Shift algorithm [20] iteratively seeks the mode of a target’s color histogram. Optical flow methods (Lucas-Kanade [21], Farnebäck [22]) estimate dense or sparse motion fields between frames. These approaches are computationally efficient but struggle with appearance changes, occlusion, and long-term identity maintenance.

section[Tracking-by-Detection Paradigm]Tracking-by-Detection Paradigm

Modern MOT predominantly follows the tracking-by-detection paradigm [23]: an independent object detector runs per frame, and a tracker associates detections across time. This decouples detection quality from tracking logic and allows plug-and-play detector upgrades.

section[SORT and Deep SORT]SORT and Deep SORT

SORT (Simple Online and Realtime Tracking) [23] combines a Kalman filter (constant velocity model) with Hungarian assignment based on IoU. It is fast and effective for low occlusion but lacks appearance information, leading to ID switches when objects cross or are occluded.

Deep SORT [24] extends SORT by integrating a deep appearance descriptor (a CNN trained on person re-identification). The association cost fuses IoU and cosine distance between embeddings. This significantly reduces ID switches but increases computational cost.

section[ByteTrack]ByteTrack

ByteTrack [1] observes that low-confidence detections (often discarded by SORT/Deep SORT) contain valuable information for tracking occluded or small objects. It introduces a two-stage association:

1. High-confidence detections ($\geq \text{trackThresh}$) matched with all tracks.
2. Low-confidence detections matched only with unmatched tracks.

This simple modification achieves state-of-the-art MOTA on MOT17 without appearance features. ByteTrack also uses a more robust Kalman filter and track lifecycle management.

section[BoT-SORT and StrongSORT]BoT-SORT and StrongSORT

BoT-SORT [25] improves ByteTrack with camera motion compensation (CMC) and a more sophisticated appearance model. StrongSORT [26] combines BoT-SORT’s improvements with an enhanced ReID model (OSNet-AIN) and adaptive appearance weighting. Both achieve higher MOTA/IDF1 at increased complexity.

section[YOLO Family]YOLO Family

- YOLOv1–v3 [12, 27, 28]: Progressive improvements in speed/accuracy.
- YOLOv4 [29]: CSPDarknet53 backbone, PANet neck, Mish activation.
- YOLOv5 [30]: PyTorch reimplementation, widely adopted, export-friendly.
- YOLOv8 [2]: Unified architecture for detection/segmentation/pose; anchor-free; improved head.

YOLOv8n (nano) is the smallest variant ($\sim 3.2\text{M}$ parameters, 12.8 MB ONNX), suitable for real-time edge/browser deployment.

section[YOLO-World: Open-Vocabulary Detection]YOLO-World: Open-Vocabulary Detection

Traditional detectors are limited to fixed training vocabularies (e.g., COCO-80). YOLO-World [3] enables *open-vocabulary* detection by conditioning on text embeddings. It uses a RepVL-PAN neck that fuses visual features with CLIP [5] text embeddings. At inference, user-supplied concepts are encoded by CLIP’s text encoder and fed as `text_features` to the detector. This enables zero-shot detection of arbitrary categories without retraining.

section[CLIP: Contrastive Language-Image Pre-training]CLIP: Contrastive Language-Image Pre-training

CLIP [5] learns joint image-text representations by contrastive pre-training on 400M image-text pairs. The text encoder (Transformer) maps tokenized captions to 512-dim embeddings; the image encoder (ViT or ResNet) maps images to the same space. Zero-shot classification uses cosine similarity between image and text embeddings. CLIP ViT-B/32 uses a Vision Transformer (ViT-B/32) image encoder and a 12-layer Transformer text encoder.

section[OSNet: Omni-Scale Network for ReID]OSNet: Omni-Scale Network for ReID

OSNet [4] is a lightweight CNN for person re-identification. Its Omni-Scale residual block captures features at multiple spatial scales simultaneously. OSNet x1.0 (~2.2M parameters) produces 512-dim embeddings and achieves strong accuracy on Market-1501 / DukeMTMC-reID. The ONNX export (Axelera mirror) expects 256x128 crops with ImageNet normalization.

section[Kalman Filter in Visual Tracking]Kalman Filter in Visual Tracking

The Kalman filter [19] is the workhorse of MOT state estimation. For 2D bounding box tracking, the state vector typically includes position, velocity, and optionally size:

$$\mathbf{x} = [x, y, v_x, v_y, w, h]^T$$

The constant velocity model assumes $x_{t+1} = x_t + v_x \Delta t$. Process noise models acceleration uncertainty; measurement noise models detection jitter. Full covariance (6x6) propagation captures correlations between state variables (e.g., position-velocity coupling).

section[Hungarian Algorithm for Data Association]Hungarian Algorithm for Data Association

The assignment problem (minimize total cost of matching tracks to detections) is solved optimally in $O(n^3)$ by the Hungarian algorithm (Kuhn-Munkres) [31, 32]. The cost matrix typically uses $1 - \text{IoU}$ or a fusion of IoU, appearance distance, and class consistency. Gating (thresholding impossible assignments) reduces the problem size.

section[Browser-Based Inference]Browser-Based Inference

ONNX Runtime Web [6] enables running ONNX models in browsers via WebAssembly (WASM), WebGL, WebGPU, and WebNN execution providers. Key milestones:

- WASM SIMD (2020): Near-native speed for CPU inference.
- WebGL (2018): GPU texture-based compute.
- WebGPU (2021): Modern GPU compute API, compute shaders, lower overhead.
- WebNN (2023): Hardware-accelerated ML abstraction layer.

Frameworks like `transformers.js` [33] demonstrate BERT, CLIP, Whisper running entirely in-browser.

section[Comparison of MOT Methods]Comparison of MOT Methods

section[Theoretical vs. Implemented]Theoretical vs. Implemented

Table 3.1: Comparison of MOT methods relevant to this work. **Bold** indicates methods implemented in TrackVision.

Method	Association	Appearance	Low-conf Dets	Camera M
SORT [23]	IoU + Hungarian			
Deep SORT [24]	IoU + App + Hungarian	(ReID)		
ByteTrack [1]	IoU + Hungarian		(2-stage)	
BoT-SORT [25]	IoU + App + Hungarian			(CMC)
StrongSORT [26]	IoU + App + Hungarian	(OSNet-AIN)		(CMC)
TrackVision (ByteTrack++)	IoU + App + Class + Center	(OSNet x1.0)	(2-stage)	

- **Implemented in TrackVision:** ByteTrack two-stage association, full 6x6 covariance Kalman, Hungarian assignment, OSNet ReID with EMA fusion, track merging, EMA smoothing, occlusion interpolation, YOLOv8n detection, YOLO-World + CLIP open-vocabulary path.
- **Literature only (not implemented):** BoT-SORT (CMC), StrongSORT (adaptive appearance), Deep SORT (original ReID), OC-SORT [34], MotionTrack [35], TransTrack [36].
- **Adapted for browser:** Kalman filter (TypeScript, Gaussian elimination), Hungarian (TypeScript, square padding), OSNet (batch-1 inference due to export constraints), CLIP tokenizer (BPE in TypeScript).

section[Key References]Key References

The following original research papers form the theoretical foundation of this work:

1. ByteTrack: Zhang et al., “ByteTrack: Multi-Object Tracking by Associating Every Detection Box,” ECCV 2022 [1].
2. YOLOv8: Jocher et al., “Ultralytics YOLOv8,” 2023 [2].
3. YOLO-World: Cheng et al., “YOLO-World: Real-Time Open-Vocabulary Object Detection,” CVPR 2024 [3].
4. OSNet: Zhou et al., “Omni-Scale Feature Learning for Person Re-Identification,” ICCV 2019 [4].
5. CLIP: Radford et al., “Learning Transferable Visual Models From Natural Language Supervision,” ICML 2021 [5].

-
6. SORT: Bewley et al., “Simple Online and Realtime Tracking,” ICIIP 2016 [23].
 7. Deep SORT: Wojke et al., “Simple Online and Realtime Tracking with a Deep Association Metric,” ICIIP 2017 [24].
 8. Kalman Filter: Kalman, “A New Approach to Linear Filtering and Prediction Problems,” 1960 [19].
 9. Hungarian Algorithm: Kuhn, “The Hungarian Method for the Assignment Problem,” 1955 [31].
 10. ONNX Runtime Web: Microsoft, “ONNX Runtime Web,” 2024 [6].

All references are formatted in IEEE style and listed in
references.bib.

chapterDataset Description / Details

section[Overview]Overview

No formal benchmark dataset is used by the current implementation of TrackVision. The system is designed for real-time inference on live camera streams or user-provided video input, not for offline evaluation on standardized MOT benchmarks (MOT17 [37], MOT20 [38], KITTI [9], TAO [39], DanceTrack [40]).

This design choice reflects the project’s primary goal: a deployable, client-side MOT application rather than a benchmark submission. Consequently, standardized quantitative metrics (MOTA, IDF1, HOTA, ID switches) cannot be computed and are not reported in this work.

section[Input Sources]Input Sources

TrackVision accepts two input modalities:

1. **Live Camera Stream:** Captured via `navigator.mediaDevices.getUserMedia` with constraints `{ video: { facingMode: 'environment', width: { ideal: 1280 }, height: { ideal: 720 } }, audio: false }`. The rear-facing camera is preferred for mobile scenarios.
2. **User-Provided Video Files:** The application supports drag-and-drop or file selection of browser-compatible video formats (MP4, WebM, MOV, etc.). The `<video>` element loads the file via `URL.createObjectURL()`.

Both sources feed into the same processing pipeline.

section[Video Characteristics]Video Characteristics

Since the input is user-dependent, the following parameters are *not fixed* but are handled dynamically by the pipeline:

section[Preprocessing Pipeline]Preprocessing Pipeline

Every frame undergoes the following deterministic preprocessing before inference:

Table 4.1: Input video parameter handling

Parameter	Handling
Resolution	Variable (camera: up to 1280x720; file: any); letterboxed to 640x640 internally
Frame Rate	Variable (typically 30 FPS); tracker assumes $\Delta t = 1/30$ s
Color Space	RGB (browser canvas); converted to CHW float32 for inference
Aspect Ratio	Preserved via letterboxing (padding); scale/offset metadata retained
Duration	Unlimited (5-minute rolling history window for storage)

1. **Capture:** `OffscreenCanvas.drawImage(video, 0, 0, 640, 640)` with letterboxing (aspect-preserving fit with padding).
2. **Metadata:** The capture returns `scale (min(640/w, 640/h))`, `offsetX`, `offsetY` (padding offsets).
3. **Planar Conversion:** `RGBA` $\rightarrow$ `float32 CHW (3 x 640 x 640)`, normalized by `1/255`.
4. **Buffer Transfer:** `ArrayBuffer` transferred to Web Workers (detached, zero-copy).

This preprocessing is implemented in `src/hooks/useOffscreenCanvas.ts` and the detection workers.

section[Object Classes]Object Classes

- **Fast Mode (YOLOv8n):** 80 fixed COCO classes (`COCO_CLASSES` in `src/workers/workerUtils.ts`). Examples: `person`, `bicycle`, `car`, `motorbike`, `bus`, `truck`, `traffic light`, `dog`, `cat`, `laptop`, `phone`, `bottle`, `cup`, `chair`, `book`, `clock`, etc.
- **Open Mode (YOLO-World):** Dynamic, user-defined concepts. The user enters comma-separated text prompts (e.g., “laptop, phone, headphones, water bottle”). Each concept is formatted as “a photo of a {concept}”, tokenized by the in-browser CLIP BPE tokenizer, and encoded by the CLIP text encoder to produce `text_features [1, C, 512]`.

section[Annotations]Annotations

No ground-truth annotations are available. The system operates in a purely *inference-only* mode. There is no training, validation, or test split. This is a fundamental limitation for quantitative MOT evaluation.

section[Test Videos Used During Development]Test Videos Used During Development

During manual development and qualitative testing, the following video sources were used:

- Live webcam (laptop integrated camera, USB webcam).
- Mobile phone camera (via dev:https LAN access).
- Sample video: one-by-one-person-detection.mp4 (Intel IoT DevKit sample, 3.3 MB, 1920x1080, 30 FPS, single person walking).

These are not benchmark datasets and were used solely for functional verification and UI testing.

section[Summary Table]Summary Table

Table 4.2: Dataset description summary

Parameter	Value
Dataset Name	<i>None (live camera / user video only)</i>
Source	Browser <code>getUserMedia()</code> / File input
Video Count	Unlimited (user-dependent)
Resolution	Variable; processed at 640x640 letterboxed
FPS	Variable; tracker $\Delta t = 1/30$ s
Object Classes	COCO-80 (fast) / User-defined (open)
Annotation Format	None (inference only)
Train/Test Split	N/A
Ground Truth	Not available

section[Implication for Evaluation]Implication for Evaluation

The absence of a benchmark dataset with ground-truth trajectories means that standardized MOT metrics cannot be computed. Chapter 9 discusses the evaluation methodology actually used (unit tests, microbenchmarks, qualitative observation) and proposes a future work path for standardized evaluation.

chapterMethodology

This chapter documents the complete implemented pipeline of TrackVision, from video acquisition to visualization. Every subsection corresponds to actual code in the repository.

section[System Overview]System Overview

TrackVision follows a tracking-by-detection architecture with four Web Workers operating in parallel, coordinated by a main-thread vision engine loop (requestAnimationFrame). The high-level data flow is:

Video → OffscreenCanvas → Detection Workers → ReID Worker → Tracker Worker → State Store → UI

Figure ?? illustrates the worker architecture.

section[Browser Video Acquisition]Browser Video Acquisition

Video input is obtained via the standard MediaDevices API:

Listing 5.1: Camera request in CameraHUD.tsx

```
1 navigator.mediaDevices.getUserMedia({
2   video: {
3     facingMode: 'environment',
4     width: { ideal: 1280 },
5     height: { ideal: 720 }
6   },
7   audio: false
8 }).then(stream => {
9   videoRef.current.srcObject = stream;
10  videoRef.current.onloadedmetadata = () => videoRef.current?.play();
11 });
```

The rear-facing camera (facingMode: 'environment') is requested by default for mobile scenarios. On desktop, the user's default camera is used. Camera permission is required; errors surface as UI banners with HTTPS guidance for insecure origins.

section[Frame Capture and Letterboxing]Frame Capture and Letterboxing

Frames are captured using OffscreenCanvas for zero-main-thread overhead:

Listing 5.2: Frame capture in useOffscreenCanvas.ts

```
1 captureFrame(targetW: number, targetH: number): CapturedFrame | null {
2   const canvas = offscreenCanvasRef.current;
3   const video = videoRef.current;
4   if (!canvas || !video || video.readyState < 2) return null;
```

```

5
6  const scale = Math.min(targetW / video.videoWidth, targetH / video.videoHeight);
7  const newW = Math.round(video.videoWidth * scale);
8  const newH = Math.round(video.videoHeight * scale);
9  const offsetX = (targetW - newW) / 2;
10 const offsetY = (targetH - newH) / 2;
11
12 const ctx = canvas.getContext('2d');
13 ctx.clearRect(0, 0, targetW, targetH);
14 ctx.drawImage(video, offsetX, offsetY, newW, newH);
15
16 const imageData = ctx.getImageData(0, 0, targetW, targetH);
17 const buffer = imageData.data.buffer.slice(0); // copy for transfer
18
19 return { buffer, data: imageData.data, width: targetW, height: targetH,
20         originalWidth: video.videoWidth, originalHeight: video.videoHeight,
21         scale, offsetX, offsetY };
22 }

```

Key properties:

- **Letterboxing:** Aspect ratio preserved; padding equally distributed.
- **Metadata:** scale, offsetX, offsetY enable mapping detection coordinates back to original video space.
- **Transfer:** ArrayBuffer is transferred (detached) to workers via postMessage, avoiding copies.

section[Preprocessing for Inference]Preprocessing for Inference

Each worker receives the transferred buffer and performs planar conversion:

Listing 5.3: Planar conversion in yoloDetectionWorker.ts

```

1 const u8 = new Uint8Array(buffer);
2 const inv255 = 1 / 255;
3 for (let i = 0, p = 0; i < numPixels; i++, p += 4) {
4   planarData[rOffset + i] = u8[p] * inv255;
5   planarData[gOffset + i] = u8[p + 1] * inv255;
6   planarData[bOffset + i] = u8[p + 2] * inv255;
7 }

```

For ReID (OSNet), crops are extracted from the letterboxed frame, resized to 256x128, and normalized with ImageNet statistics (mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]).

section[Object Detection]Object Detection

Two detection paths are implemented:

subsection[Fast Mode: YOLOv8n]Fast Mode: YOLOv8n

- **Model:** yolov8n.onnx (12.8 MB, COCO-80)

- **Worker:** `yoloDetectionWorker.ts`
- **Input:** `float32[1, 3, 640, 640]`
- **Output:** Transposed `[1, 84, 8400]` ($84 = 4 \text{ bbox} + 80 \text{ classes}$) or non-transposed
- **Postprocessing:** Sigmoid calibration, class-aware NMS (IoU 0.45), coordinate mapping via `scale/offsetX/offsetY`

subsection[Open-Vocabulary Mode: YOLO-World +
CLIP]Open-Vocabulary Mode: YOLO-World + CLIP

- **Model:** `yoloworld.onnx` (418 MB) + `clip_text_encoder.onnx` (254 MB)
- **Worker:** `yoloWorldWorker.ts`
- **CLIP Tokenizer:** `clipTokenizer.ts` (BPE, `vocab/merges` from `public/models/clip-tokenizer/`)
- **Text Encoding:** “a photo of a {concept}” $\rightarrow$ `tokenize(int64[77])` $\rightarrow$ CLIP encoder (WASM) $\rightarrow$ L2-normalize $\rightarrow$ `text_features[1, C, 512]`
- **YOLO-World Input:** `images[1, 3, 640, 640]` + `text_features[1, C, 512]`
- **Output:** Post-sigmoid scores `[1, N, C]` + XYXY boxes `[1, N, 4]`
- **Postprocessing:** `postprocessWorld` decodes scores/boxes in letterbox space
- **Lazy Loading:** CLIP loads on first concept set; embeddings computed once per concept set

section[Detection Filtering and Class-Aware NMS]Detection
Filtering and Class-Aware NMS

After detection, class-aware NMS runs on the main thread
(`applyNMSWithClass` in `workerUtils.ts`):

1. Group detections by `classId`.
2. Within each class, sort by confidence descending.
3. Iterate: keep highest, suppress others with $\text{IoU} > 0.45$.
4. Merge results across classes.

This prevents cross-class suppression (e.g., a high-confidence “person” does not suppress a nearby “laptop”).

section[ReID Feature Extraction]ReID Feature Extraction

ReID is conditional – only active when a valid OSNet model is loaded (`reidReady && !reidIsFallback`).

1. Crop Mapping: Video-space boxes $\rightarrow$ letterbox coordinates using `scale/offsetX/offsetY`.
2. Batch-1 Inference: Each crop processed sequentially (OSNet export rejects stacked batches).
3. Preprocessing: Resize 256x128, ImageNet mean/std, RGB $\rightarrow$ CHW.
4. Inference: `extractBatch` in `reidWorker.ts` returns 512-dim L2-normalized embedding per crop.
5. Attachment: Non-zero embeddings attached to detections if length matches.

If OSNet is unavailable, the tracker falls back to $\text{IoU} + \text{center distance} + \text{class consistency}$ only.

section[Track Prediction]Track Prediction

Before association, every active track predicts its next state via Kalman filter:

Listing 5.4: Track prediction in `tracker.ts`

```
1 predict(): BBox {
2   const predicted = this.kalman.predict();
3   this.bbox = predicted;
4   return predicted;
5 }
```

The Kalman filter (`kalman.ts`) maintains a 6x6 covariance matrix with constant acceleration process noise.

section[Data Association]Data Association

The core of ByteTrack++ is the four-stage association

(`ByteTracker.update` in `tracker.ts`):

subsection[Stage 1: High-Confidence Association]Stage 1:

High-Confidence Association Match detections with $\text{score} \geq \text{trackThresh}$ (capped at 0.4) against all existing tracks using Hungarian algorithm on combined cost matrix.

subsection[Stage 2: Low-Confidence Association]Stage 2:

Low-Confidence Association Match detections with $0.1 < \text{score} < \text{trackThresh}$ against only the *unmatched* tracks from Stage 1.

subsection[Stage 3: New Track Creation]Stage 3: New Track Creation Unmatched high-confidence detections spawn new tracks (TENTATIVE state, new Kalman filter).

subsection[Stage 4: Missed Track Handling]Stage 4: Missed Track Handling Unmatched tracks increment `timeSinceUpdate`; if > 5 frames $\rightarrow$ LOST; if > 30 frames (`maxTimeLost`) $\rightarrow$ REMOVED.

subsection[Cost Matrix Composition]Cost Matrix Composition

For each track-detection pair, the combined cost is:


$$\text{cost} = \begin{cases} 0.6 \cdot (1 - \text{IoU}) + 0.3 \cdot (1 - \text{cosine_sim}) + 0.1 \cdot \text{class_cost} + 0.1 \cdot \text{center_cost} \\ 0.8 \cdot (1 - \text{IoU}) + 0.15 \cdot \text{center_cost} + 0.05 \cdot \text{class_cost} \end{cases}$$

Where:

- `class_cost` = 0 if same class, 0.4 otherwise
- `center_cost` = $\min(1, \frac{\text{center_dist}}{2 \cdot \max(\text{dims})})$
- Gating: costs $> 1 - \text{matchThresh}$ (default 0.3) set to ∞ before Hungarian

section[Track Lifecycle Management]Track Lifecycle Management

Tracks transition through states:



- confirmedHits: detections with score ≥ 0.7
- Recovery: LOST $\rightarrow$ TRACKED if re-associated within maxTimeLost
- Label stabilization: Majority vote over 10 class labels
- Score smoothing: EMA $\alpha = 0.3$ (score $\leftarrow 0.7 \cdot \text{old} + 0.3 \cdot \text{new}$)

section[Track Merging (Split Identity Recovery)]Track Merging (Split Identity Recovery)

If two confirmed tracks of the *same class* have IoU > 0.5 and appearance similarity > 0.7 , the older track absorbs the younger (EMA embedding fusion, younger marked REMOVED).

section[EMA Smoothing and Interpolation]EMA Smoothing and Interpolation

- BBox EMA: $\alpha = 0.7$ on bbox coordinates.
- Center EMA: $\alpha = 0.8$ on center points.
- Occlusion Interpolation: Up to 5 frames of Kalman-predicted positions blended with smoothed bbox when timeSinceUpdate > 0 .

section[Output Rendering]Output Rendering

The main thread renders tracking overlays on a canvas
(drawTrackingHUD in lib/draw.ts):

- Bounding boxes with track IDs, class labels, confidence.
- Ghost mode: Fading trajectory trails (15 frames, opacity decay).

- Follow mode: Camera locked to selected track, dimmed background.
- Scene Map: 640x480 top-down density map with live track dots.
- Timeline: Gantt-chart horizontal bars per track ID (5-min window).

section[Telemetry Collection]Telemetry Collection

Per-frame latency breakdown collected in `useVisionEngine.ts`:

- `preprocessMs`, `inferenceMs`, `postprocessMs`, `trackingMs`, `renderMs`, `frameMs`
- `executionProvider` (WebNN/WebGPU/WebGL/WASM)
- `fps` (rolling 500ms window)
- History: 60-sample rolling buffers for latency/FPS charts

section[Five-Minute Rolling History]Five-Minute Rolling History

The Zustand store (`store.ts`) maintains a sliding window:

Listing 5.5: History pruning in `store.ts`

```
1 const cutoffTime = frame.timestamp - state.maxHistoryMs; // 5 * 60 * 1000
2 while (newFrames.length > 0 && newFrames[0].timestamp < cutoffTime) {
3   newFrames.shift();
4 }
```

Track metadata (trajectories, distance, speed, bearing) is updated per frame and pruned to 300 history points per track.

section[Replay Mode]Replay Mode

When `isReplay=true`, the live loop pauses. `currentTime` scrubs through stored `frames[]`, rendering historical frames with full track state reconstruction.

chapterAlgorithm / Flowchart

This chapter presents the core algorithms implemented in TrackVision: the Kalman filter, Hungarian assignment, ByteTrack++ association pipeline, and ReID fusion. All algorithms are transcribed directly from the TypeScript implementation.

section[Kalman Filter (Full 6x6 Covariance)]Kalman Filter (Full 6x6 Covariance)

The Kalman filter maintains a state vector $\mathbf{x} \in \mathbb{R}^6$ and covariance matrix $\mathbf{P} \in \mathbb{R}^{6 \times 6}$.

subsection[State Definition]State Definition

$$\mathbf{x} = \begin{bmatrix} x & y & v_x & v_y & w & h \end{bmatrix}^T$$

where (x, y) is the box center, (v_x, v_y) velocity, (w, h) width and height.

subsection[State Transition (Predict)]State Transition (Predict)

$$\mathbf{F} = \begin{bmatrix} 1 & 0 & \Delta t & 0 & 0 & 0 \\ 0 & 1 & 0 & \Delta t & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}, \quad \mathbf{x}_{k|k-1} = \mathbf{F}\mathbf{x}_{k-1|k-1}$$

Process noise $\mathbf{Q}$ models constant acceleration uncertainty:

$$\mathbf{Q} = q \begin{bmatrix} \frac{\Delta t^4}{4} & 0 & \frac{\Delta t^3}{2} & 0 & 0 & 0 \\ 0 & \frac{\Delta t^4}{4} & 0 & \frac{\Delta t^3}{2} & 0 & 0 \\ \frac{\Delta t^3}{2} & 0 & \Delta t^2 & 0 & 0 & 0 \\ 0 & \frac{\Delta t^3}{2} & 0 & \Delta t^2 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0.01 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0.01 \end{bmatrix}$$

$$\mathbf{P}_{k|k-1} = \mathbf{F}\mathbf{P}_{k-1|k-1}\mathbf{F}^T + \mathbf{Q}$$

Velocity is clamped to ± 1000 px/frame to prevent runaway predictions.

subsection[Measurement Update]Measurement Update
 Measurement vector $\mathbf{z} = [m_x, m_y, 0, 0, m_w, m_h]^T$ (position + size; velocity not measured).

Measurement matrix $\mathbf{H}$ maps state to measurement space:

$$\mathbf{H} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

Measurement noise $\mathbf{R}$ (diagonal):

$$\mathbf{R} = \text{diag}(r, r, 1, 1, r, r), \quad r = \text{measurementNoise}$$

$$\text{Innovation: } \mathbf{y} = \mathbf{z} - \mathbf{H}\mathbf{x}_{k|k-1}$$

$$\text{Innovation covariance: } \mathbf{S} = \mathbf{H}\mathbf{P}_{k|k-1}\mathbf{H}^T + \mathbf{R}$$

$$\text{Kalman gain: } \mathbf{K} = \mathbf{P}_{k|k-1}\mathbf{H}^T\mathbf{S}^{-1}$$

$$\text{State update: } \mathbf{x}_{k|k} = \mathbf{x}_{k|k-1} + \mathbf{K}\mathbf{y}$$

$$\text{Covariance update: } \mathbf{P}_{k|k} = (\mathbf{I} - \mathbf{K}\mathbf{H})\mathbf{P}_{k|k-1}$$

subsection[Matrix Inversion]Matrix Inversion $\mathbf{S}^{-1}$ is computed via Gaussian elimination (6x6) in `KalmanFilter.matrixInverse()`. If singular, identity is returned as fallback.

section[Hungarian Algorithm (Kuhn-Munkres)]Hungarian Algorithm (Kuhn-Munkres)

Solves the assignment problem: minimize $\sum_{i,j} C_{ij}x_{ij}$ subject to each track/detection assigned at most once.

Algorithm 1 Hungarian Algorithm (Square Cost Matrix)

```

procedure HUNGARIAN( $C \in \mathbb{R}^{n \times n}$ )
   $u, v \leftarrow \mathbf{0}_{n+1}$  ▷ dual variables
   $p, way \leftarrow \mathbf{0}_{n+1}$ 
  for  $i = 1 \dots n$  do
     $p[0] \leftarrow i, minv \leftarrow \infty_{n+1}, used \leftarrow false_{n+1}, j_0 \leftarrow 0$ 
    repeat
       $used[j_0] \leftarrow true$ 
       $i_0 \leftarrow p[j_0], \delta \leftarrow \infty, j_1 \leftarrow 0$ 
      for  $j = 1 \dots n$  do
        if  $\neg used[j]$  then
           $cur \leftarrow C[i_0 - 1][j - 1] - u[i_0] - v[j]$ 
          if  $cur < minv[j]$  then  $minv[j] \leftarrow cur, way[j] \leftarrow j_0$ 
          if  $minv[j] < \delta$  then  $\delta \leftarrow minv[j], j_1 \leftarrow j$ 

      for  $j = 0 \dots n$  do
        if  $used[j]$  then  $u[p[j]] \leftarrow u[p[j]] + \delta, v[j] \leftarrow v[j] - \delta$ 
        else  $minv[j] \leftarrow minv[j] - \delta$ 

       $j_0 \leftarrow j_1$ 
       $p[j_0] = 0$ 
      repeat
         $j_1 \leftarrow way[j_0], p[j_0] \leftarrow p[j_1], j_0 \leftarrow j_1$ 
      until  $j_0 = 0$ 

    return  $\{(p[j] - 1, j - 1, C[p[j] - 1][j - 1]) \mid j = 1..n, p[j] \neq 0\}$ 

```

subsection[Algorithm Outline]Algorithm Outline

subsection[Gating and Rectangular Matrices]Gating and Rectangular Matrices For $n_{tracks} \neq n_{dets}$, the cost matrix is padded to square with cost 10^6 (infinite). Assignments with cost $> \text{maxCost}$ are discarded.

section[ByteTrack++ Association Pipeline]ByteTrack++ Association Pipeline

The `ByteTracker.update(detections)` method implements four stages:

section[Cost Matrix Composition]Cost Matrix Composition

For each track i and detection j :

Algorithm 2 ByteTrack++ Update (Per Frame)

```

procedure UPDATE(detections)
  for all track  $\in$  tracks do                                      $\triangleright$  Predict all
    track.predict()
    high  $\leftarrow \{d \in \text{detections} \mid d.\text{score} \geq \text{trackThresh}\}$ 
    low  $\leftarrow \{d \in \text{detections} \mid 0.1 < d.\text{score} < \text{trackThresh}\}$ 
                                                                     $\triangleright$  Stage 1: High-confidence association
    result1  $\leftarrow$  Hungarian(tracks, high, matchThresh)
    for all (t, d)  $\in$  result1.matched do
      tracks[t].update(high[d])
                                                                     $\triangleright$  Stage 2: Low-confidence association
    remaining  $\leftarrow$  result1.unmatchedTracks
    result2  $\leftarrow$  Hungarian(tracks[remaining], low, matchThresh)
    for all (t, d)  $\in$  result2.matched do
      tracks[remaining[t]].update(low[d])
                                                                     $\triangleright$  Stage 3: New tracks from unmatched high-conf
    for all d  $\in$  result1.unmatchedDetections do
      tracks.push(new Track(high[d]))
                                                                     $\triangleright$  Stage 4: Mark missed
    for all t  $\in$  result2.unmatchedTracks do
      tracks[remaining[t]].markMissed()
    MergeSimilarTracks()
                                                                     $\triangleright$  Split identity recovery
    return tracks.filter(¬deleted)
=0

```

$$\text{iou} = \text{IoU}(\text{predBBBox}_i, \text{detBBBox}_j) \quad (6.1)$$

$$\text{iouCost} = 1 - \text{iou} \quad (6.2)$$

$$\text{classCost} = \begin{cases} 0 & \text{if } \text{class}_i = \text{class}_j \\ 0.4 & \text{otherwise} \end{cases} \quad (6.3)$$

$$\text{centerDist} = \|\text{center}(\text{predBBBox}_i) - \text{center}(\text{detBBBox}_j)\|_2 \quad (6.4)$$

$$\text{centerCost} = \min\left(1, \frac{\text{centerDist}}{2 \cdot \max(\text{dims})}\right) \quad (6.5)$$

If both have embeddings:

$$\text{appCost} = 1 - \max(0, \text{cosineSim}(\text{emb}_i, \text{emb}_j)) \quad (6.6)$$

$$\text{combinedCost} = 0.6 \cdot \text{iouCost} + 0.3 \cdot \text{appCost} + 0.1 \cdot \text{classCost} + 0.1 \cdot \text{centerCost} \quad (6.7)$$

Otherwise (no embeddings):

$$\text{combinedCost} = 0.8 \cdot \text{iouCost} + 0.15 \cdot \text{centerCost} + 0.05 \cdot \text{classCost}$$

Gating: if $\text{combinedCost} > 1 - \text{matchThresh}$ (default 0.3), set to 10^6 .

section[ReID Embedding Fusion]ReID Embedding Fusion


When a track with embedding e_{old} receives a detection with embedding e_{new} :

$$e_{fused} = \text{normalize}((1 - \alpha)e_{old} + \alpha e_{new}), \quad \alpha = 0.3$$

Normalization: $e \leftarrow e / \|e\|_2$ (L2 projection onto unit hypersphere).

During track merging (same class, $\text{IoU} > 0.5$, $\text{app sim} > 0.7$), the older track absorbs the younger with the same EMA fusion.

section[Track Lifecycle State Machine]Track Lifecycle State Machine



- **hits:** total associated detections
- **confirmedHits:** associated detections with $\text{score} \geq 0.7$
- **timeSinceUpdate:** frames since last successful association
- **maxTimeLost = 30** (configurable)

section[Flowchart: End-to-End Frame Processing]Flowchart: End-to-End Frame Processing

6.1 Computational Complexity

Table 6.1: Per-frame computational complexity (asymptotic)

Component	Complexity
YOLOv8n Inference	$O(1)$ (fixed 640x640 input, model-dependent)
Class-Aware NMS	$O(N \log N)$ per class (N = detections)
OSNet ReID (per crop)	$O(1)$ (fixed 256x128, sequential batch-1)
Kalman Predict/Update	$O(1)$ per track (6x6 matrix ops)
Hungarian Assignment	$O((T + D)^3)$ (T tracks, D detections)
Track Merging	$O(T^2)$ pairwise IoU
EMA Smoothing	$O(T)$

Practical bottleneck: **Model inference** (YOLO/OSNet/CLIP) dominates frame time. JS-side tracking pipeline: **0.31–1.04 ms/frame** at 5–20 detections (Vitest benchmark, Node 22).

CHAPTER 7

Source Code

This chapter presents curated excerpts from the TrackVision implementation. Each listing is a direct extract from the repository with explanatory context. The complete source is available at <https://github.com/zubershk/TrackVision>.

section[Kalman Filter: Predict Step]Kalman Filter: Predict Step
src/lib/kalman.ts – Core prediction with constant acceleration process noise.

Listing 7.1: KalmanFilter.predict()

```
1 predict(): BBox {
2   const F = [
3     [1, 0, this.dt, 0, 0, 0],
4     [0, 1, 0, this.dt, 0, 0],
5     [0, 0, 1, 0, 0, 0],
6     [0, 0, 0, 1, 0, 0],
7     [0, 0, 0, 0, 1, 0],
8     [0, 0, 0, 0, 0, 1]
9   ];
10
11   const q = this.processNoise;
12   const dt = this.dt;
13   const Q = [
14     [q * dt**4 / 4, 0, q * dt**3 / 2, 0, 0, 0],
15     [0, q * dt**4 / 4, 0, q * dt**3 / 2, 0, 0],
16     [q * dt**3 / 2, 0, q * dt**2, 0, 0, 0],
17     [0, q * dt**3 / 2, 0, q * dt**2, 0, 0],
18     [0, 0, 0, 0, q * 0.01, 0],
19     [0, 0, 0, 0, 0, q * 0.01]
20   ];
21
22   // x = F * x
23   const newState = F.map(row =>
24     row.reduce((sum, val, i) => sum + val * Object.values(this.state)[i], 0)
25   );
26
27   this.state = { /* velocity clamped to +/-1000 */ };
28
29   // P = F * P * F^T + Q
30   const FP = this.matrixMultiply(F, this.covariance);
31   const FTP = this.matrixTranspose(FP);
32   const FPFT = this.matrixMultiply(FP, FTP);
33   this.covariance = this.matrixAdd(FPFT, Q);
34
35   return this.stateToBBox();
36 }
```

section[Kalman Filter: Update Step]Kalman Filter: Update Step src/lib/kalman.ts – Gaussian elimination for 6x6 covariance inversion.

Listing 7.2: KalmanFilter.update()

```

1 update(measurement: BBox): BBox {
2   const [x, y, w, h] = measurement;
3   const mx = x + w/2, my = y + h/2, mw = w, mh = h;
4
5   const H = [ /* measurement matrix */ ];
6   const R = [ /* measurement noise */ ];
7
8   const z = [mx, my, 0, 0, mw, mh];
9   const Hx = H.map(row => row.reduce((sum, val, i) => sum + val * Object.values(this.state)[i], 0));
10  const innovation = z.map((val, i) => val - Hx[i]);
11
12  const HP = this.matrixMultiply(H, this.covariance);
13  const HT = this.matrixTranspose(H);
14  const HPHT = this.matrixMultiply(HP, HT);
15  const S = this.matrixAdd(HPHT, R);
16
17  const PHT = this.matrixMultiply(this.covariance, HT);
18  const Sinv = this.matrixInverse(S) || this.identity(6);
19  const K = this.matrixMultiply(PHT, Sinv);
20
21  const Ky = K.map(row => row.reduce((sum, val, i) => sum + val * innovation[i], 0));
22  // state update, covariance update...
23  return this.stateToBBox();
24 }

```

section[Track Update: EMA Smoothing and Label Stabilization]Track Update: EMA Smoothing and Label Stabilization src/lib/tracker.ts – Track.update() with EMA smoothing, label voting, and embedding fusion.

Listing 7.3: Track.update()

```

1 update(detection: Detection): void {
2   this.kalman.update(detection.bbox);
3   this.bbox = this.kalman.stateToBBox();
4
5   // BBox EMA (alpha=0.7)
6   this.smoothedBBox = [
7     this.smoothedBBox[0] * this.bboxEMAAlpha + this.bbox[0] * (1 - this.bboxEMAAlpha),
8     this.smoothedBBox[1] * this.bboxEMAAlpha + this.bbox[1] * (1 - this.bboxEMAAlpha),
9     this.smoothedBBox[2] * this.bboxEMAAlpha + this.bbox[2] * (1 - this.bboxEMAAlpha),
10    this.smoothedBBox[3] * this.bboxEMAAlpha + this.bbox[3] * (1 - this.bboxEMAAlpha),
11  ] as BBox;
12
13  // Score EMA (alpha=0.3)
14  this.score = this.score * 0.7 + detection.score * 0.3;
15
16  // Label stabilization: majority vote over last 10 labels
17  this.labelHistory.push(detection.class);
18  if (this.labelHistory.length > 10) this.labelHistory.shift();
19
20  const counts = new Map<string, number>();
21  let maxCount = 0, stableLabel = this.className;
22  for (const label of this.labelHistory) {
23    const c = (counts.get(label) || 0) + 1;
24    counts.set(label, c);
25    if (c > maxCount) { maxCount = c; stableLabel = label; }
26  }
27  this.className = stableLabel;
28
29  // Embedding fusion (alpha=0.3) with L2 re-normalization
30  if (detection.embedding) {

```

```

31   if (this.embedding) {
32     this.embedding = this.updateEmbedding(this.embedding, detection.embedding, 0.3);
33   } else {
34     this.embedding = detection.embedding;
35   }
36   this.lastEmbeddingUpdate = Date.now();
37 }
38
39 // Lifecycle transitions
40 if (this.state === TrackState.TENTATIVE && (this.hits >= 3 || this.confirmedHits >= 2)) {
41   this.state = TrackState.CONFIRMED;
42   this.status = TrackStatus.TRACKED;
43 }
44 }

```

section[ByteTracker: Four-Stage Association]ByteTracker: Four-Stage Association src/lib/tracker.ts – Core ByteTracker.update() implementing the four stages.

Listing 7.4: ByteTracker.update()

```

1 update(detections: Detection[]): Track[] {
2   for (const track of this.tracks) track.predict();
3
4   const highScores = detections.filter(d => d.score >= this.trackThresh);
5   const lowScores = detections.filter(d => d.score > 0.1 && d.score < this.trackThresh);
6
7   const unmatchedTracks = [...this.tracks];
8   const newTracks: Track[] = [];
9
10  // Stage 1: High-score detections <-> all tracks
11  const highResult = this.matchDetections(unmatchedTracks, highScores, this.matchThresh);
12  for (const m of highResult.matched) unmatchedTracks[m.trackIdx].update(highScores[m.detIdx]);
13
14  // Stage 2: Low-score detections <-> remaining tracks
15  const remaining1 = highResult.unmatchedTracks.map(i => unmatchedTracks[i]);
16  const lowResult = this.matchDetections(remaining1, lowScores, this.matchThresh);
17  for (const m of lowResult.matched) {
18    const origIdx = highResult.unmatchedTracks[m.trackIdx];
19    unmatchedTracks[origIdx].update(lowScores[m.detIdx]);
20  }
21
22  // Stage 3: New tracks from unmatched high-score detections
23  for (const detIdx of highResult.unmatchedDetections) {
24    const det = highScores[detIdx];
25    if (det) newTracks.push(new Track(this.nextId++, det, new KalmanFilter(det.bbox)));
26  }
27
28  // Stage 4: Mark unmatched tracks as missed
29  for (const subIdx of lowResult.unmatchedTracks) {
30    const origIdx = highResult.unmatchedTracks[subIdx];
31    unmatchedTracks[origIdx]?.markMissed();
32  }
33
34  this.mergeSimilarTracks();
35  // Collect and filter...
36  return this.tracks.filter(t => t.status !== TrackStatus.REMOVED);
37 }

```

section[Cost Matrix: IoU + Appearance + Class + Center]Cost Matrix: IoU + Appearance + Class + Center src/lib/tracker.ts – Combined cost computation in matchDetections().

Listing 7.5: matchDetections() cost computation

```

1 private matchDetections(tracks: Track[], detections: Detection[], maxCost: number): MatchResult {

```

```

2  const costMatrix: number[][] = [];
3
4  for (let i = 0; i < tracks.length; i++) {
5      const row: number[] = [];
6      const predBBBox = tracks[i].getPredictedBBBox();
7
8      for (let j = 0; j < detections.length; j++) {
9          const iou = getIoU(predBBBox, detections[j].bbox);
10         const iouCost = 1 - iou;
11
12         const hasEmb = Boolean(tracks[i].embedding && detections[j].embedding);
13         let appCost = 0;
14         if (hasEmb) {
15             const sim = tracks[i].cosineSimilarity(detections[j].embedding);
16             appCost = 1 - Math.max(0, sim);
17         }
18
19         const classCost = tracks[i].className === detections[j].class ? 0 : 0.4;
20         const centerDist = this.getCenterDistance(predBBBox, detections[j].bbox);
21         const maxDim = Math.max(predBBBox[2], predBBBox[3], detections[j].bbox[2], detections[j].bbox[3], 1);
22         const centerCost = Math.min(1, centerDist / (maxDim * 2));
23
24         let combinedCost: number;
25         if (hasEmb) {
26             combinedCost = iouCost * this.iouWeight // 0.6
27                 + appCost * this.embeddingWeight // 0.3
28                 + classCost * this.classWeight // 0.1
29                 + centerCost * 0.1;
30         } else {
31             combinedCost = iouCost * 0.8 + centerCost * 0.15 + classCost * 0.05;
32         }
33         row.push(combinedCost);
34     }
35     costMatrix.push(row);
36 }
37 return this.useHungarian ? hungarianMatch(costMatrix, maxCost) : { ... };
38 }

```

section[ReID Batch Extraction]ReID Batch Extraction

src/workers/reidWorker.ts – Sequential batch-1 inference for OSNet.

Listing 7.6: ReID extractBatch()

```

1  async function extractBatch(imageData: ImageData, bboxes: BBox[]): Promise<(Float32Array | null)[]> {
2      if (!modelLoaded || !session) return bboxes.map(() => null);
3
4      const embeddings: (Float32Array | null)[] = [];
5      for (let b = 0; b < bboxes.length; b++) {
6          try {
7              const inputData = cropImage(imageData, bboxes[b], INPUT_SIZE, OSNET_NORMALIZE);
8              const inputTensor = new ort.Tensor('float32', inputData, [1, 3, 256, 128]);
9              const results = await session.run({ [inputName]: inputTensor });
10             const output = results[outputName].data as Float32Array;
11
12             const embedding = new Float32Array(512);
13             embedding.set(output.slice(0, 512));
14             // L2 normalize
15             let norm = 0;
16             for (let i = 0; i < 512; i++) norm += embedding[i] * embedding[i];
17             norm = Math.sqrt(norm);
18             if (norm > 0) for (let i = 0; i < 512; i++) embedding[i] /= norm;
19
20             embeddings.push(embedding);
21         } catch { embeddings.push(null); }
22     }
23     return embeddings;
24 }

```

section[Class-Aware NMS]Class-Aware NMS

src/workers/workerUtils.ts – Per-class suppression preventing cross-class suppression.

Listing 7.7: applyNMSWithClass()

```

1 export function applyNMSWithClass(detections: Detection[], iouThreshold: number): Detection[] {
2   const byClass = new Map<number, Detection[]>();
3   for (const det of detections) {
4     if (!byClass.has(det.classId)) byClass.set(det.classId, []);
5     byClass.get(det.classId)!.push(det);
6   }
7
8   const results: Detection[] = [];
9   for (const [, dets] of byClass) {
10    dets.sort((a, b) => b.score - a.score);
11    for (let i = 0; i < dets.length; i++) {
12      if (dets[i].score <= 0) continue;
13      results.push(dets[i]);
14      for (let j = i + 1; j < dets.length; j++) {
15        if (getIoU(dets[i].bbox, dets[j].bbox) > iouThreshold) {
16          dets[j].score = 0;
17        }
18      }
19    }
20  }
21  return results;
22 }

```

section[Vision Engine Main Loop]Vision Engine Main Loop

src/hooks/useVisionEngine.ts – Per-frame requestAnimationFrame loop coordinating all workers.

Listing 7.8: useVisionEngine processFrame()

```

1 const processFrame = useCallback(async () => {
2   if (!state.isTracking || state.isReplay || !videoRef.current || !trackerInitializedRef.current) return;
3   if (state.visionMode === 'fast' && !yoloReady) return;
4   if (state.visionMode === 'open' && !yoloWorldReady) return;
5   if (isDetectingRef.current) return;
6
7   const frame = captureFrame(640, 640);
8   if (!frame) return scheduleNext();
9
10  isDetectingRef.current = true;
11  const frameStart = performance.now();
12
13  // Clone frame for ReID before buffer transfer
14  const imageDataForReID = (reidReady && !reidIsFallback)
15    ? new ImageData(frame.data.slice(), frame.width, frame.height)
16    : null;
17
18  // 1. DETECT
19  let detections = state.visionMode === 'fast'
20    ? (await yoloDetect(frame, confidenceThreshold)).results
21    : (await yoloWorldDetect(frame, confidenceThreshold)).results;
22
23  detections = applyNMSWithClass(detections, 0.45);
24
25  // 2. ReID (if available)
26  if (detections.length > 0 && imageDataForReID) {
27    const boxes = detections.map(d => [
28      d.bbox[0] * frame.scale + frame.offsetX,
29      d.bbox[1] * frame.scale + frame.offsetY,
30      d.bbox[2] * frame.scale,
31      d.bbox[3] * frame.scale
32    ]);
33    const embeddings = await extractBatch(imageDataForReID, boxes);
34    for (let i = 0; i < detections.length; i++) {
35      if (embeddings[i]?.length && hasNonZeroValues(embeddings[i]!))

```

```

36     detections[i].embedding = embeddings[i]!;
37   }
38 }
39
40 // 3. TRACK
41 const activeTracks = (await updateTracker(detections)) ?? [];
42
43 // 4. STORE + TELEMETRY
44 addFrameData({ timestamp, tracks: tracksData }, activeTracksInfo, {
45   fps, inferenceMs, preprocessMs, postprocessMs, trackingMs, frameMs,
46   executionProvider, deviceAcceleration
47 });
48
49 } finally { isDetectingRef.current = false; scheduleNext(); }
50 }, [/* deps */]);

```

section[Hungarian Algorithm (Core)]Hungarian Algorithm (Core) src/lib/matching.ts – Square matrix assignment with dual variables.

Listing 7.9: hungarianAlgorithm()

```

1 function hungarianAlgorithm(costMatrix: number[][], maxCost: number): { row: number; col: number; cost: number }[] {
2   const n = costMatrix.length;
3   const u = new Array(n + 1).fill(0);
4   const v = new Array(n + 1).fill(0);
5   const p = new Array(n + 1).fill(0);
6   const way = new Array(n + 1).fill(0);
7
8   for (let i = 1; i <= n; i++) {
9     p[0] = i;
10    const minv = new Array(n + 1).fill(Infinity);
11    const used = new Array(n + 1).fill(false);
12    let j0 = 0;
13
14    do {
15      used[j0] = true;
16      const i0 = p[j0];
17      let delta = Infinity, j1 = 0;
18      for (let j = 1; j <= n; j++) if (!used[j]) {
19        const cur = costMatrix[i0 - 1][j - 1] - u[i0] - v[j];
20        if (cur < minv[j]) { minv[j] = cur; way[j] = j0; }
21        if (minv[j] < delta) { delta = minv[j]; j1 = j; }
22      }
23      for (let j = 0; j <= n; j++) {
24        if (used[j]) { u[p[j]] += delta; v[j] -= delta; }
25        else { minv[j] -= delta; }
26      }
27      j0 = j1;
28    } while (p[j0] !== 0);
29
30    do { const j1 = way[j0]; p[j0] = p[j1]; j0 = j1; } while (j0 !== 0);
31  }
32
33  const assignment = [];
34  for (let j = 1; j <= n; j++) if (p[j] !== 0 && p[j] <= costMatrix.length && j <= costMatrix[0].length) {
35    const cost = costMatrix[p[j] - 1][j - 1];
36    if (cost <= maxCost) assignment.push({ row: p[j] - 1, col: j - 1, cost });
37  }
38  return assignment;
39 }

```

chapterOutput Screenshots

This section presents visual output from the TrackVision application. Screenshots were captured from the running production build (`npm run preview`) using Playwright automation at <http://localhost:4173>.

Important Disclosure: The TrackVision application acquires video exclusively via `getUserMedia()` (live camera). In the headless automation environment used for documentation generation, camera access is unavailable. Therefore, screenshots showing actual detection/tracking output (bounding boxes, tracking IDs, trajectories) could not be captured. These figures are marked as placeholders and must be manually captured by running the application locally with a camera.

section[Application Interface]Application Interface
section[Tracking Output (Placeholders – Require Live Camera)]Tracking Output (Placeholders – Require Live Camera)
section[Telemetry Visualization]Telemetry Visualization
section[Summary]Summary

Table 8.1: Screenshot capture summary

Fig	Description	Status
1	Landing Page	• Real capture
2	Command Center (Standby)	• Real capture
3	Command Center (Ready)	• Real capture
4	Vision Panel	* Placeholder
5	Analytics Panel	* Placeholder
6	Track Inspector	* Placeholder
7	Scene Map	* Placeholder
8	Timeline	* Placeholder
9	Live View	* Placeholder
10	Command Palette	• Real capture
11	Mobile View	• Real capture
12	Final Composite	• Real capture
13–20	Tracking Output	* Placeholders (camera required)

To capture placeholders: Run `npm run dev` (or `npm run dev:https` for mobile), grant camera permission, and use OS screenshot tools or implement video file upload support.

TrackVision v2.0
CLIENT-SIDE CV & TRACKING

FeaturesArchitecturePlaygroundBenchmarksUse CasesCode

Launch Engine →

WebGPU Accelerated • 100% Client-Side • Zero Data Uploads

Precision Object Intelligence. Running Entirely in Your Browser.

Real-time multi-object tracking with **ByteTrack++** temporal persistence, dual neural detection (**YOLOv8** + **YOLO-World**), and **OSNet** appearance re-identification — powered locally via WebGPU with zero cloud latency.

Launch Command Center →Explore Architecture

100% LOCAL
Zero Server Ingestion
100% private in-browser compute

60+ FPS LATENCY
15-30ms Per Frame
Hardware WebGPU / WASM SIMD

BYTETRACK++
6x6 Covariance Kalman
High & low score dual association

OSNET RE-ID
512-dim Embeddings
Occlusion identity recovery

LIVE VISION SIMULATORByteTrack++ • Interactive 60 FPS

FPS: 3LATENCY: 17msTRACKS: 4

LOCKED: #01 person 98%
VELOCITY: 15.0 px/s
BEARING: 200°
RE-ID HASH: 0a0f36...4821

#03 drone 92%

#01 person 98%

#02 laptop 95%

#04 vehicle 99%

Click anywhere to ping sonar & spawn target • Click entity to lock focus

Pause+ Spawn Entity

LAYERS: Ghost Trails (G) • Kalman Vectors (V) • Radar Sweep •

Fast (COCO 80)Open-Vocabulary

BUILT ON OPEN NEURAL & CLIENT-SIDE WEB STANDARDS

Native
Runtime

ONNX Runtime Web
Neural Backend

ByteTrack++
Temporal Tracking

YOLOv8 & YOLO-World
Zero-Shot Vision

OSNet v1.0
512-dim Re-ID

WASM SIMD
Vector fallback

React 19 & Typescript
Frontend Stack

CORE CAPABILITIES

Engineered for Precision & Performance

Every layer of TrackVision is architected in pure TypeScript and ONNX Runtime Web for native client-side execution with zero compromise on accuracy.

PRIVACY FIRST

100% Client-Side Intelligence

Zero video frames or sensitive telemetry leave your device. All neural networks execute locally in your browser memory via WebGPU and WASM SIMD.

CLOUD TRANSMISSION0 BYTES (SECURE)

DUAL ARCHITECTURE

YOLOv8 & Open-Vocabulary Vision

Seamlessly toggle between ultra-fast COCO 80-class object detection and zero-shot YOLO-World + CLIP open-vocabulary text-conditioned prompt detection.

INFERENCE SPEED60+ FPS (16ms)

STATE ESTIMATION

ByteTrack++ & 6x6 Kalman Filter

Custom TypeScript implementation with a full 6x6 covariance matrix modeling position, velocity, and aspect ratio under constant acceleration.

KALMAN STATE[x, y, vx, vy, w, h]




Figure 8.2: Command Center – Initial state after clicking “Launch Command Center”. Shows top navigation (Live, Tracks, Vision, History, Scene, Analytics), center viewport with “Start Tracking” button, and Engine Thinking Stream pill. Models not yet initialized.




Figure 8.3: Command Center – Ready state after model initialization (YOLOv8n loaded). “Start Tracking” button enabled, telemetry stream active, Live badge shows standby FPS.

[INSERT SCREENSHOT: Vision Panel – Detection Engine Selector]

Figure 8.4: Vision Panel (Detection Engine) – Fast Mode (YOLOv8n, COCO-80) vs Open Mode (YOLO-World + CLIP) selector. Open mode shows concepts input textarea with comma-separated prompts. Runtime fallback status shows execution provider.

[INSERT SCREENSHOT: Analytics Panel – Real-time Telemetry]

Figure 8.5: Analytics Panel (Settings tab) – Real-time telemetry cards (FPS, frame time, inference latency, tracking latency), object volume AreaChart, class distribution BarChart. Updates live during tracking.

[INSERT SCREENSHOT: Track Inspector – Active Tracks List]

Figure 8.6: Track Inspector (Tracks tab) – Active tracks list with track ID, class label, confidence score, status (NEW/TRACKED/LOST), and selection controls. Click to enable Follow mode.

[INSERT SCREENSHOT: Scene Map – Top-Down Density]

Figure 8.7: Scene Map (Scene tab) – 640x480 top-down spatial density map with live track position dots, color-coded by track ID. Updates in real-time during tracking.

[INSERT SCREENSHOT: Timeline – 5-Minute History Scrubber]

Figure 8.8: Timeline (History tab) – Gantt-chart style horizontal bars per track ID, 5-minute sliding window, variable playback speeds (0.5x, 1x, 2x, 4x), scrubber for frame-perfect seek.




Figure 8.9: Command Palette – Activated via Ctrl+K. Quick actions: Start/Stop Tracking, Toggle Ghost Mode, Toggle Follow Mode, Clear History, Switch Vision Mode, Open Settings.




Figure 8.10: Mobile Responsive View (375x667) – Bottom navigation bar (Live, Tracks, History, Scene, Vision, Analytics), stacked panels, touch-optimized controls.




Figure 8.11: Full Desktop Command Center – Live view with camera placeholder, left sidebar (Tracks panel shown), top navigation with telemetry (FPS, latency), Engine Thinking Stream, corner controls (Ghost Mode, Follow Mode), bottom live HUD.

[INSERT SCREENSHOT: Detection Result – Bounding Boxes + Labels + Confidence]

Figure 8.12: Early detection frame – YOLOv8n output showing bounding boxes, COCO class labels, and confidence scores. No tracking IDs assigned yet (first frame).

[INSERT SCREENSHOT: Multi-Object Tracking – Persistent IDs]

Figure 8.13: Three or more objects tracked simultaneously with persistent integer IDs, class labels, and confidence scores. IDs remain stable across frames.

[INSERT SCREENSHOT: Persistent Identity – Same Object Across Frames]

Figure 8.14: Identity continuity demonstration – same physical object maintains the same track ID across multiple frames despite motion and minor appearance changes.

[INSERT SCREENSHOT: Ghost Mode – Fading Trajectory Trails]

Figure 8.15: Ghost mode enabled – 15-frame fading trajectory trails with opacity decay, showing historical path of each tracked object.

[INSERT SCREENSHOT: Follow Mode – Target Lock]

Figure 8.16: Follow mode enabled – camera viewport locked to selected track, background dimmed, selected track highlighted with persistent ID and telemetry.

[INSERT SCREENSHOT: Occlusion Handling – Track Maintenance]

Figure 8.17: Object temporarily occluded – track maintains ID via Kalman prediction (interpolation frames), reappears with same ID after occlusion.

[INSERT SCREENSHOT: Crowded Scene – ID Separation]

Figure 8.18: Multiple overlapping objects in close proximity – distinct IDs maintained, track merging activates only when IoU > 0.5, same class, and appearance similarity > 0.7.

[INSERT SCREENSHOT: ReID Appearance Match – Re-acquisition]

Figure 8.19: Track re-acquired after prolonged occlusion – OSNet appearance embedding fusion enables correct ID assignment despite spatial gap. Requires OSNet model loaded.




Figure 8.20: Real-time Telemetry – FPS history (line chart), latency breakdown (stacked area: preprocess, inference, postprocess, tracking, render), object count per frame, class distribution bar chart. Updates every frame during tracking.

chapterResult & Accuracy

section[Evaluation Methodology]Evaluation Methodology

TrackVision’s tracking performance was evaluated through the following evidence-based methods:

1. **Implementation-Level Unit Tests:** 49 Vitest tests validating Kalman filtering, Hungarian assignment, ByteTrack lifecycle, appearance fusion, class-aware NMS, and store history pruning.
2. **Runtime Microbenchmarks:** JavaScript-side pipeline timing (NMS + ByteTrack association + Kalman predict/update) measured over 1000 frames in Node.js 22.
3. **Telemetry Collection:** Per-frame latency breakdown (preprocess, inference, postprocess, tracking, render) and FPS recorded during manual qualitative testing.
4. **Qualitative Observation:** Manual verification of identity continuity, occlusion handling, track merging, and responsiveness using live camera and sample video.
5. **Session Initialization Timing:** Model loading and ONNX session creation times measured on Windows 11 / Chrome with WebGPU.
6. **Bundle Size Analysis:** Production build artifact sizes measured from dist/.

Critical Limitation: The current implementation does not compute standardized MOT benchmark metrics such as MOTA, IDF1, HOTA, ID switches, false positives, or false negatives. No formal benchmark dataset (MOT17, MOT20, KITTI, TAO, DanceTrack) is used, and no ground-truth annotations are available. The evaluation therefore focuses on engineering validation and computational performance rather than benchmark-oriented MOT accuracy.

section[Implementation Validation: Unit Tests]Implementation Validation: Unit Tests

All 49 unit tests pass (npm run test):

Table 9.1: Unit test coverage by module

Module	Tests
Kalman Filter (<code>kalman.test.ts</code>)	9
Hungarian/Greedy Matching (<code>matching.test.ts</code>)	8
ByteTrack Lifecycle + Appearance (<code>tracker.test.ts</code>)	20
Worker Utils / NMS (<code>workerUtils.test.ts</code>)	7
Zustand Store History (<code>store.test.ts</code>)	5
Total	49

- **Kalman:** Predict/update correctness, covariance propagation, matrix inversion, velocity clamping, state reset.
- **Matching:** Hungarian optimal assignment, gating, rectangular padding, greedy fallback, edge cases (empty inputs).
- **ByteTrack:** TENTATIVE→CONFIRMED→LOST→REMOVED transitions, confirmation thresholds ($\text{hits} \geq 3$ or $\text{confirmedHits} \geq 2$), occlusion recovery (re-association after 3 missed frames), ID persistence during smooth motion, separate IDs for spatially distinct objects, low-score detection handling, label stabilization via majority vote.
- **Appearance Fusion:** EMA blending ($\alpha=0.3$) with L2 re-normalization, crossing disambiguation when IoU ambiguous (appearance-dominant tuning).
- **Store:** 5-minute rolling window frame pruning, track metadata persistence, event logging.

section[Pipeline Microbenchmarks]Pipeline Microbenchmarks

Table 9.2: Tracking pipeline latency (JS only, no model inference)

Detections/Frame	ms/frame
5	0.31
10	0.51
20	1.04




Figure 9.1: Tracking pipeline latency scales sub-linearly with detection count. Model inference (not shown) dominates end-to-end frame time.

JavaScript-side tracking pipeline (NMS + ByteTrack association + Kalman predict/update) measured over 1000 frames per configuration in Node.js 22 (Vitest benchmark):

Interpretation: The JS tracking overhead is negligible (< 1.1 ms/frame even at 20 detections). End-to-end frame time is dominated by model inference (YOLO/OSNet/CLIP), which is hardware-dependent (WebGPU vs. WASM).

section[Session Initialization Times]Session Initialization Times
Measured on Windows 11, Chrome (WebGPU execution provider), production build:

Table 9.3: ONNX session initialization time (model load + compile)

Model	Size	Session Init
YOLOv8n (detection)	12.8 MB	~1.8 s
OSNet x1.0 (ReID)	8.8 MB	~1.9 s
YOLO-World (open)	418 MB	~3.7 s

Notes:

- Initialization includes model download (if not cached), ONNX parsing, graph optimization, and execution provider compilation.
- WebGPU provider adds shader compilation time; WASM is faster to initialize but slower at inference.
- YOLO-World’s large size (418 MB) dominates its initialization time.
- CLIP text encoder (~ 254 MB) loads lazily on first concept set in open mode.

section[Bundle Sizes (Production Build)]Bundle Sizes (Production Build)

Measured from `dist/assets/` after `npm run build` (raw, not gzipped):

Table 9.4: Production bundle sizes

Chunk	Size
Main bundle (<code>index.js</code>)	345 KB
Charts chunk (<code>charts.js</code> , lazy)	373 KB
Vendor chunk	36 KB
YOLO Detection Worker	404 KB
YOLO-World Worker	404 KB
ReID Worker	400 KB
Tracker Worker	10 KB
CSS	81 KB
ONNX Runtime WASM (<code>ort-wasm-simd-threaded.jsep.wasm</code>)	26.2 MB

The WASM binary is cached by the browser and loaded on-demand by workers. Total initial download (without optional models): ~ 1.7 MB (JS + CSS + WASM).

section[Qualitative Tracking Behavior]Qualitative Tracking Behavior

Manual testing with live camera and sample video
(`one-by-one-person-detection.mp4`) observed:

- Identity Continuity: Tracks maintain stable IDs during smooth motion; re-association succeeds after brief occlusion (≤ 5 frames).

- **Occlusion Handling:** Kalman prediction interpolates up to 5 frames during occlusion; track status transitions TRACKED → LOST after 6 missed frames.
- **Track Merging:** When two tracks of same class overlap ($\text{IoU} > 0.5$) with high appearance similarity (> 0.7), older track absorbs younger.
- **Label Stabilization:** Majority vote over 10-frame history corrects transient misclassifications.
- **Open-Vocabulary Detection:** YOLO-World + CLIP successfully detects user-specified concepts (e.g., “water bottle”, “headphones”) when full models are downloaded.
- **ReID Fusion:** With OSNet loaded, appearance embeddings improve crossing disambiguation; without OSNet, tracker gracefully falls back to IoU + center + class.

section[Standardized MOT Evaluation – Future Work]Standardized MOT Evaluation – Future Work

The following standardized evaluation is not currently implemented but is recommended for future work:

Table 9.5: Recommended MOT metrics for future benchmark evaluation

Metric	Description
MOTA	Multiple Object Tracking Accuracy (detection + association)
IDF1	ID F1-score (identity consistency)
HOTA	Higher Order Tracking Accuracy (unified detection+association)
ID Switches	Count of identity changes
FP / FN	False positives / false negatives
MOTP	Multiple Object Tracking Precision (localization)

Required for standardized evaluation:

1. Annotated benchmark dataset (MOT17 train/val, MOT20, or DanceTrack).
2. Ground-truth trajectory parser (MOT format: `<frame>`, `<id>`, `<bb_left>`, `<bb_top>`, `<bb_width>`, `<bb_height>`, `<conf>`, `<class>`, `<vis>`).

3. Evaluation pipeline computing MOTA, IDF1, HOTA per [41, 42].
4. Deterministic inference seed for reproducibility.
5. CI integration for regression tracking.

section[Summary of Achieved Results]Summary of Achieved Results

Table 9.6: Summary of measured results (achieved)

Metric	Value
Unit tests passing	49 / 49
Pipeline latency (5 dets)	0.31 ms/frame
Pipeline latency (10 dets)	0.51 ms/frame
Pipeline latency (20 dets)	1.04 ms/frame
YOLOv8n init (WebGPU)	~1.8 s
OSNet init (WebGPU)	~1.9 s
YOLO-World init (WebGPU)	~3.7 s
Main bundle size	345 KB
Total JS/CSS (excl. WASM)	~1.7 MB
ONNX WASM binary	26.2 MB
TypeScript strict check	Pass (<code>tsc -noEmit</code>)

Explicitly NOT achieved (by design):

- MOTA / IDF1 / HOTA scores
- ID switch count
- False positive / negative rates
- Detection mAP on COCO
- End-to-end FPS claims (device-dependent; model inference dominates)

section[Conclusion on Accuracy]Conclusion on Accuracy

TrackVision demonstrates correct algorithmic implementation (validated by 49 unit tests) and low computational overhead (sub-millisecond JS tracking pipeline). The system achieves real-time performance on consumer hardware with WebGPU acceleration. However, no quantitative MOT accuracy claim can be made without benchmark dataset evaluation. The qualitative behavior confirms the tracker functions as designed: IDs persist, occlusions are handled via Kalman prediction, and appearance fusion (when available) improves crossing disambiguation.

chapterApplications

section[Demonstrated Application]Demonstrated Application

Real-time, privacy-preserving, browser-based multiple object tracking using live camera or user-provided video, with zero server dependencies, zero video upload, and zero installation.

This is the only application fully demonstrated by the current implementation. The TrackVision Command Center provides:

- Live detection and tracking with persistent IDs
- Ghost mode (trajectory visualization)
- Follow mode (target lock)
- Time Machine (5-minute history replay)
- Scene Map (spatial density)
- Analytics (volume, class distribution, telemetry)
- Dual detection modes: Fast (COCO-80) and Open (user-defined concepts)
- Responsive UI (desktop + mobile)
- PWA installation for offline access

section[Potential Applications]Potential Applications

The following applications are technically plausible given TrackVision's capabilities but have not been experimentally validated in this project. They are presented as future deployment directions.

subsection[Traffic Monitoring]Traffic Monitoring

- Vehicle counting, classification (car, truck, bus, motorcycle), speed estimation via trajectory analysis.
- Intersection monitoring, queue length estimation, wrong-way detection.

- **Advantage: Runs on edge devices (traffic cameras with browser) without cloud connectivity.**

subsection[Pedestrian Monitoring]Pedestrian Monitoring
 subsection[Surveillance and Security]Surveillance and Security
 subsection[Sports Analytics]Sports Analytics
 subsection[Retail Analytics]Retail Analytics
 subsection[Crowd Analysis]Crowd Analysis
 subsection[Autonomous Systems and Robotics]Autonomous
 Systems and Robotics
 subsection[Smart City Monitoring]Smart City Monitoring
 subsection[Industrial Monitoring]Industrial Monitoring
 section[Deployment Advantages for Applications]Deployment
 Advantages for Applications

Table 10.1: TrackVision deployment advantages for real-world applications

Requirement	TrackVision Property
Privacy / GDPR	Zero video upload; fully client-side
Latency	No network round-trip; local inference
Hardware	Runs on any WebGPU/WASM device (phone, laptop, edge)
Installation	Zero (browser) or PWA (offline)
Scalability	Static hosting (CDN); no server compute
Cost	Near-zero (static hosting only)
Cross-platform	Chrome, Edge, Firefox, Safari (desktop + mobile)
Customization	Open-vocabulary concepts without retraining

section[Limitations for Production Deployment]Limitations for
 Production Deployment

Before deploying in any of the above domains, the following must
 be addressed:

1. **Benchmark Validation:** MOTA/IDF1/HOTA on domain-specific data (traffic, retail, etc.).
2. **Robustness:** Extended testing under varying lighting, weather, occlusion density, camera angles.
3. **Video File Support:** Current implementation only supports live camera; file upload needed for recorded footage analysis.

4. **Export/Logging:** Structured output (JSON, CSV, MOT format) for downstream analytics.
5. **Multi-Camera:** Single-camera only; multi-camera fusion requires additional work.
6. **Alerting/Integration:** Webhooks, MQTT, or API for real-time event notification.

section[Conclusion]Conclusion

TrackVision’s core contribution is demonstrating that a complete, production-quality MOT stack can run entirely in-browser. This architectural approach (client-side, zero-server, privacy-first) is directly applicable to any domain requiring real-time object tracking where privacy, latency, cost, or connectivity constraints preclude cloud-based solutions. The potential applications listed above are credible deployment targets *pending* the validation steps outlined in Section [10](#).

chapterConclusion

section[Summary of Work]Summary of Work

This project presented TRACKVISION, a fully client-side, real-time multiple object tracking system implemented in TypeScript/React that runs entirely within a web browser using ONNX Runtime Web. The system achieves the following:

1. **Dual Detection Modes:** Fast path with YOLOv8n (COCO-80, 12.8 MB) and open-vocabulary path with YOLO-World + CLIP ViT-B/32 (in-browser text encoding).
2. **Custom ByteTrack++ Tracker:** Full 6x6 covariance Kalman filter, Hungarian data association, four-stage ByteTrack association, track lifecycle management (TENTATIVE→CONFIRMED→LOST→OSNet x1.0 ReID with EMA fusion ($\alpha=0.3$), track merging, and EMA smoothing.
3. **Browser-Native Pipeline:** OffscreenCanvas frame capture at 640x640 letterboxed, zero-copy buffer transfer to four Web Workers (YOLO, YOLO-World, ReID, Tracker), class-aware NMS on main thread.
4. **Comprehensive UI:** Command Center with Live/Ghost/Follow/Replay/Scout/Map/Analytics/Tracks tabs, 5-minute rolling history, real-time telemetry (FPS, latency breakdown, execution provider), Command Palette, responsive design.
5. **Validation:** 49 passing Vitest unit tests covering all core algorithms; pipeline microbenchmarks (0.31–1.04 ms/frame JS tracking overhead); session initialization measurements; bundle size analysis.
6. **Deployment Ready:** Static hosting (Netlify/Vercel/Cloudflare Pages) with COOP/COEP headers, Docker + nginx, HTTPS dev server for mobile camera access, PWA with Service Worker.

section[What Worked Well]What Worked Well

- **Web Worker Architecture:** Clean separation of concerns; typed message protocols; no main-thread blocking during inference.
- **Hardware Acceleration Cascade:** WebNN $\rightarrow$ WebGPU $\rightarrow$ WebGL $\rightarrow$ WASM fallback works transparently; users get best available performance.
- **Open-Vocabulary Path:** In-browser CLIP BPE tokenizer + text encoder eliminates external embedding services; concepts encoded once per set.
- **Tracker Robustness:** ByteTrack++ handles occlusion (Kalman interpolation), ID switches (appearance fusion), and split identities (track merging) correctly per unit tests.
- **Telemetry & History:** 5-minute rolling window with frame-perfect replay enables post-hoc analysis.
- **TypeScript Strict Mode:** Zero any without justification; compile-time safety across workers, hooks, stores, and components.

section[Major Limitations]Major Limitations

1. **No Benchmark Dataset / MOT Metrics:** The system does not compute MOTA, IDF1, HOTA, or ID switches. No formal evaluation on MOT17, MOT20, KITTI, etc. Quantitative accuracy claims cannot be made.
2. **Camera-Only Input:** No video file upload support; limits testing reproducibility and recorded footage analysis.
3. **Model Download Dependency:** Open mode (418 MB + 254 MB) and ReID (8.8 MB) require explicit download; placeholder stubs gate these features.
4. **Device/Browser Variability:** WebGPU availability inconsistent (Chrome/Edge best; Firefox behind flag; Safari limited); WASM fallback significantly slower.
5. **No Multi-Camera / 3D:** Single monocular camera only; no stereo, no 3D localization, no camera motion compensation.

6. No Automated E2E Tests: Only unit tests; no Playwright/Cypress integration tests for UI flows.
7. No Video Export: Live canvas only; no MP4/WebM recording of tracked output.

section[Technical Contributions]Technical Contributions

1. Complete in-browser MOT stack: First known implementation combining YOLOv8, YOLO-World, CLIP, ByteTrack, Kalman, Hungarian, OSNet, and ONNX Runtime Web in a single client-side application.
2. In-browser CLIP text encoding: Real BPE tokenizer (TypeScript) + ONNX CLIP text encoder (WASM) for open-vocabulary detection without external APIs.
3. ByteTrack++ in TypeScript: Full covariance Kalman, two-stage association, appearance fusion, track merging, EMA smoothing – all type-safe and tested.
4. Execution Provider Abstraction: `createAcceleratedSession` cascade (`WebNN`→`WebGPU`→`WebGL`→`WASM`) reusable across workers.
5. Production Deployment Model: Static hosting + COOP/COEP + PWA + Docker + HTTPS dev workflow.

section[Future Work]Future Work

1. Standardized MOT Evaluation: Integrate MOT17/MOT20/DanceTrack datasets; implement MOTA/IDF1/HOTA computation; add to CI.
2. Video File Upload: Support `<input type="file">` for recorded footage analysis and reproducible testing.
3. Automated E2E Testing: Playwright tests for critical user flows (start tracking, switch modes, replay, export).

4. **Video Export:** MediaRecorder API to save tracked output as MP4/WebM.
5. **Multi-Camera Fusion:** Synchronized multi-view tracking with triangulation.
6. **Camera Motion Compensation:** Implement CMC (as in BoT-SORT) for moving camera scenarios.
7. **Model Optimization:** INT8 quantization for smaller/faster WASM/WebGPU inference.
8. **Structured Output:** JSON/CSV/MOT-format logging for downstream analytics.

section[Final Statement]Final Statement

TrackVision demonstrates that modern web technologies (WebGPU, WASM, WebNN, OffscreenCanvas, Web Workers) are sufficient to deploy a production-grade multiple object tracking system entirely on the client side. This architecture eliminates server costs, preserves user privacy, enables offline operation, and scales to millions of users via static CDN hosting. While standardized benchmark evaluation remains future work, the implementation is algorithmically sound (49 passing tests), computationally efficient (< 1.1 ms JS tracking overhead), and deployment-ready.

TrackVision – Precision Object Intelligence

chapterBibliography

- [1] Y. Zhang, P. Sun, Y. Jiang, D. Yu, Z. Yuan, P. Luo, W. Liu, and X. Wang, “Bytetrack: Multi-object tracking by associating every detection box,” *European Conference on Computer Vision (ECCV)*, pp. 1–21, 2022.
- [2] G. Jocher, A. Chaurasia, and J. Qiu, “Ultralytics yolov8,” *GitHub Repository*, 2023.
- [3] T. Cheng, X. Wang, L. Huang, and W. Liu, “Yolo-world: Real-time open-vocabulary object detection,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 16900–16910, 2024.
- [4] K. Zhou, Y. Yang, A. Cavallaro, and T. Xiang, “Omni-scale feature learning for person re-identification,” *IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 3701–3711, 2019.
- [5] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning transferable visual models from natural language supervision,” *International Conference on Machine Learning (ICML)*, vol. 139, pp. 8748–8763, 2021.
- [6] Microsoft, “Onnx runtime web,” *GitHub Repository*, 2024.
- [7] O. Russakovsky, J. Deng, H. Su, J. Krause, S. Satheesh, S. Ma, Z. Huang, A. Karpathy, A. Khosla, M. Bernstein, A. C. Berg, and L. Fei-Fei, “Imagenet large scale visual recognition challenge,” *International Journal of Computer Vision*, vol. 115, no. 3, pp. 211–252, 2015.
- [8] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, “Microsoft coco: Common objects in context,” *European Conference on Computer Vision (ECCV)*, pp. 740–755, 2014.

- [9] A. Geiger, P. Lenz, C. Stiller, and R. Urtasun, “Vision meets robotics: The kitti dataset,” *International Journal of Robotics Research*, vol. 32, no. 11, pp. 1231–1237, 2013.
- [10] R. Girshick, J. Donahue, T. Darrell, and J. Malik, “Rich feature hierarchies for accurate object detection and semantic segmentation,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 580–587, 2014.
- [11] S. Ren, K. He, R. Girshick, and J. Sun, “Faster r-cnn: Towards real-time object detection with region proposal networks,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 28, pp. 91–99, 2015.
- [12] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, “You only look once: Unified, real-time object detection,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 779–788, 2016.
- [13] W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. Reed, C.-Y. Fu, and A. C. Berg, “Ssd: Single shot multibox detector,” *European Conference on Computer Vision (ECCV)*, pp. 21–37, 2016.
- [14] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, “Focal loss for dense object detection,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2980–2988, 2017.
- [15] D. S. Bolme, J. R. Beveridge, B. A. Draper, and Y. M. Lui, “Visual object tracking using adaptive correlation filters,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2544–2550, 2010.
- [16] A. Lukežič, T. Vojř, L. Čehovin Zajc, J. Matas, and M. Kristan, “Discriminative correlation filter with channel and spatial reliability,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 6309–6318, 2017.

- [17] L. Bertinetto, J. Valmadre, J. F. Henriques, A. Vedaldi, and P. H. S. Torr, “Fully-convolutional siamese networks for object tracking,” *European Conference on Computer Vision (ECCV)*, pp. 850–865, 2016.
- [18] B. Li, J. Yan, W. Wu, Z. Zhu, and X. Hu, “High performance visual tracking with siamese region proposal network,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 8971–8980, 2018.
- [19] R. E. Kalman, “A new approach to linear filtering and prediction problems,” *Journal of Basic Engineering*, vol. 82, no. 1, pp. 35–45, 1960.
- [20] D. Comaniciu, V. Ramesh, and P. Meer, “Kernel-based object tracking,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 25, no. 5, pp. 564–577, 2003.
- [21] B. D. Lucas and T. Kanade, “An iterative image registration technique with an application to stereo vision,” *Proceedings of the 7th International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 674–679, 1981.
- [22] G. Farnebäck, “Two-frame motion estimation based on polynomial expansion,” *Scandinavian Conference on Image Analysis (SCIA)*, pp. 363–370, 2003.
- [23] A. Bewley, Z. Ge, L. Ott, F. Ramos, and B. Upcroft, “Simple online and realtime tracking,” *IEEE International Conference on Image Processing (ICIP)*, pp. 3464–3468, 2016.
- [24] N. Wojke, A. Bewley, and D. Paulus, “Simple online and real-time tracking with a deep association metric,” *IEEE International Conference on Image Processing (ICIP)*, pp. 3645–3649, 2017.
- [25] N. Aharon, R. Orfaig, and B.-Z. Bobrov, “Bot-sort: Robust associations multi-pedestrian tracking,” *arXiv preprint arXiv:2206.14651*, 2022.

- [26] Y. Du, Z. Zhao, Y. Song, Y. Zhao, F. Su, T. Gong, and H. Meng, “Strongsort: Make deepsort great again,” *IEEE Transactions on Multimedia*, vol. 25, pp. 8725–8737, 2023.
- [27] J. Redmon and A. Farhadi, “Yolo9000: Better, faster, stronger,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 7263–7271, 2017.
- [28] J. Redmon and A. Farhadi, “Yolov3: An incremental improvement,” *arXiv preprint arXiv:1804.02767*, 2018.
- [29] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao, “Yolov4: Optimal speed and accuracy of object detection,” *arXiv preprint arXiv:2004.10934*, 2020.
- [30] G. Jocher *et al.*, “Yolov5,” *GitHub Repository*, 2021.
- [31] H. W. Kuhn, “The hungarian method for the assignment problem,” *Naval Research Logistics Quarterly*, vol. 2, no. 1-2, pp. 83–97, 1955.
- [32] J. Munkres, “Algorithms for the assignment and transportation problems,” *Journal of the Society for Industrial and Applied Mathematics*, vol. 5, no. 1, pp. 32–38, 1957.
- [33] Xenova, “Transformers.js: State-of-the-art machine learning for the web,” *GitHub Repository*, 2023.
- [34] J. Cao, X. Weng, R. Khirodkar, J. Pang, and K. Kitani, “Observation-centric sort: Rethinking sort for robust multi-object tracking,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 9686–9696, 2023.
- [35] L. Chen, H. Ai, Z. Zhuang, and C. Shang, “Motiontrack: Learning robust short-term and long-term motions for multi-object tracking,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [36] P. Sun, J. Cao, Y. Jiang, Z. Yuan, S. Bai, K. Kitani, and P. Luo, “Transtrack: Multiple object tracking with transformer,” *arXiv preprint arXiv:2012.15460*, 2020.

- [37] A. Milan, L. Leal-Taixé, I. Reid, S. Roth, and K. Schindler, “Mot16: A benchmark for multi-object tracking,” *arXiv preprint arXiv:1603.00831*, 2016.
- [38] P. Dendorfer, H. Rezatofighi, A. Milan, J. Shi, S. Roth, K. Schindler, and I. Reid, “Mot20: A benchmark for multi-object tracking in crowded scenes,” *arXiv preprint arXiv:2003.09003*, 2020.
- [39] A. Dave, T. Khurana, P. Tokmakov, C. Schmid, and D. Ramanan, “Tao: A large-scale benchmark for tracking any object,” *European Conference on Computer Vision (ECCV)*, pp. 436–454, 2020.
- [40] P. Sun, J. Cao, Y. Jiang, R. Zhang, Z. Yuan, S. Bai, K. Kitani, and P. Luo, “Dancetrack: Multi-object tracking in uniform appearance and diverse motion,” *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 20993–21002, 2022.
- [41] K. Bernardin and R. Stiefelhagen, “Evaluating multiple object tracking performance: The clear mot metrics,” *EURASIP Journal on Image and Video Processing*, vol. 2008, no. 1, pp. 1–10, 2008.
- [42] J. Luiten, A. Osep, P. Dendorfer, P. Torr, A. Geiger, L. Leal-Taixé, and B. Leibe, “Hota: A higher order metric for evaluating multi-object tracking,” *International Journal of Computer Vision*, vol. 129, pp. 548–578, 2021.